

# Welcome Advanced Software Architecture

Hello



## About me...



- 25+ years in the industry
- 20+ years in teaching
- Certified Cloud architect
- Passionate about learning
- Also, passionate about  
Reese's Cups!

# Prerequisites

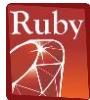
## This course assumes you:

- Have hands-on coding experience in a modern programming language (C#, C++, Java, Python, etc.)
- Have experience building moderate-sized software applications

# Why study these subjects?

- Building software systems is “easy” – building software systems that perform, scale, are testable, and that provide lasting value is not
- Architectural quality does not happen by accident
- Having said that, there are some key principles and practices that help us be successful
- Spending some time laying a foundation in those principles and practices is time well spent

# We teach over 400 technology topics



# You experience our impact on a daily basis!



# My pledge to you

I will...

- Make this interactive
- Ask you questions
- Ensure everyone can speak
- Use an on-screen timer

# Objectives

**At the end of this course, you will be able to:**

- Gain an in-depth understanding of key software architecture concepts and methodologies
- Grasp how to refactor complex systems to achieve test coverage and scalability
- Recognize the tradeoffs between alternate designs depending upon staffing, time to market, and scalability constraints

# Agenda (Some of the Topics We'll Be Covering)

- Overview of key technologies, practices, and strategies that inform and enable software architecture
- Software architecture styles
- Solution & evolutionary architecture
- Architecting for the Cloud
- Software design patterns

# Origin of Design Patterns

# First Use of “Design Pattern”



- Term first coined by an architect and anthropologist – Christopher Alexander
- Presented a new language construct based around an entity called a “pattern”
- Pattern describes a problem and provides a reusable (and proven) solution to that problem

# Gang of Four (GoF)



- Initial search for “Gang of Four” on [www.wikipedia.org](https://www.wikipedia.org)

The screenshot shows a search results page for "gang of four" on Wikipedia. The search bar at the top contains the query "gang of four". Below the search bar, there are several search results listed:

- Gang of Four**  
Chinese political faction
- Gang of Four (band)**  
English rock band
- Gang of Four (SDP)**  
Breakaway group from UK Labour party in 1981
- Gang of Four (pro-Contra)**
- Gang of Four (film)**  
1988 film
- Gang of Four (Harlem)**

On the right side of the search results, there are links to other Wikimedia projects:

- Wikivoyage: Free travel
- Wikinews: Free news
- Wikiquote: Free quote
- Wikisource
- Wikispecies

Not exactly what  
we're looking for...

# Gang of Four (GoF)



- Search for “Gang of Four Design Patterns”

Special page

## Search results

Q Gang of Four Design Patterns

Advanced search: Sort by relevance X

Search in: Article X

The page "Gang of Four Design Patterns" does not exist. You can [create a draft and submit it for review](#), but consider checking the search results below to see whether the topic is already covered.

**Design Patterns**  
Design Patterns: Elements of Reusable Object-Oriented Software (1994) is a software engineering book describing software **design patterns**. The book was...  
15 KB (1,711 words) - 10:31, 3 August 2022

**Singleton pattern**  
pattern is a software **design pattern** that restricts the instantiation of a class to a singular instance. One of the well-known "Gang of Four" design patterns...  
9 KB (830 words) - 04:13, 19 October 2022

**Builder pattern**  
Builder **design pattern** is to separate the construction of a complex object from its representation. It is one of the **Gang of Four design patterns**. The Builder...

That's more like it...

# Gang of Four (GoF)



- Group of 4 authors who wrote the book titled “Design Patterns: Elements of Reusable Object-Oriented Software” (1994)
- Erich Gamma, Richard Helm, Ralph Johnson, and John Vlissides
- Includes detail on 3 types of patterns

# Gang of Four (GoF)



Creational

Structural

Behavioral

# Gang of Four (GoF)



Supports creation of objects indirectly  
(in a more loosely-coupled fashion);  
enables association of logic to  
determine what and how to create

Creational

Structural

Behavioral

# Gang of Four (GoF)



Creational

Structural

Behavioral

About class and object composition; using inheritance and extension to build out entity hierarchies that match with the “real world” and enable layering in new functionality in an architecturally sound manner

# Gang of Four (GoF)



Creational

Structural

Behavioral

Mainly manages concepts of communication between objects – building out a messaging system that allows us to break a larger problem into smaller pieces but still coordinate

# Software Architecture Patterns vs. Design Patterns



Architecture Style

Architecture Pattern

Design Pattern

# Software Architecture Patterns vs. Design Patterns



Describe the macro structure of a system (e.g., event-driven architecture)

Architecture Style

Architecture Pattern

Design Pattern

# Software Architecture Patterns vs. Design Patterns



Architecture Style

Architecture Pattern

Reusable structural pattern that can be used to enable the architecture style (e.g., CQRS)

Design Pattern

# Software Architecture Patterns vs. Design Patterns



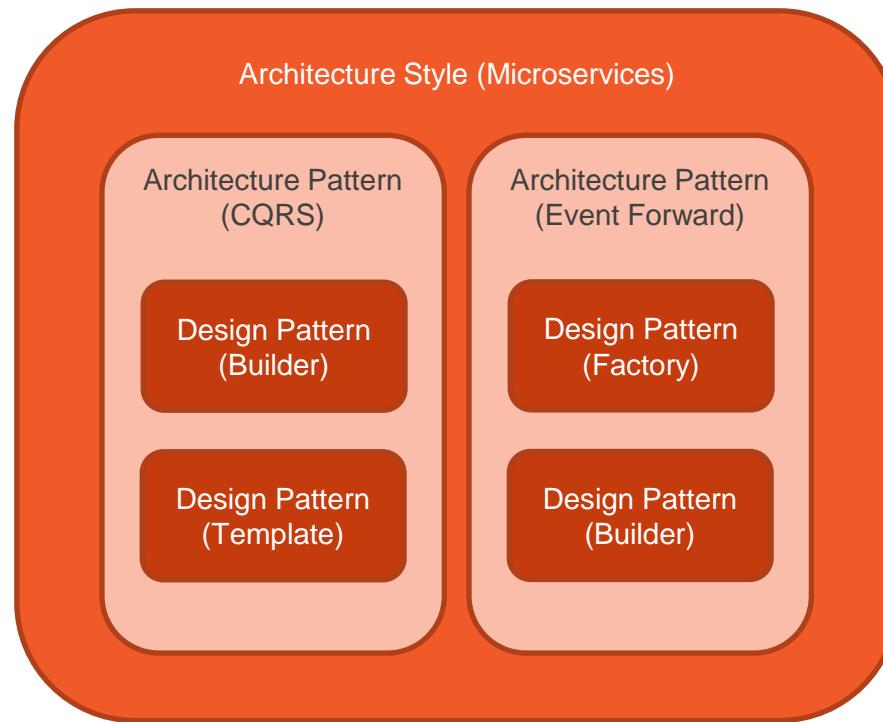
Architecture Style

Architecture Pattern

Design Pattern

Impacts how source code is designed  
and structured

# Architecture Style, Architecture Pattern, Design Pattern



## QUICK CHECK

Architecture Patterns &  
Design Patterns

For the following “boxes”, specify whether you believe the concept represented by the box is an architecture pattern or a design pattern (given the distinction/definitions just discussed).

Layered

Microkernel

Singleton

Observer

Factory

Adapter

Microservices

Chain of  
Responsibility

# High-Level Overview of Concepts that Enable Software Architecture

# Waterfall vs. Iterative Development

# Waterfall Software Development

A sequential development process that flows like a waterfall through all phases of a project (requirements, design, implementation, testing, and deployment for example), with each phase completely wrapping up before the next phase begins.

Key aspects:

- Majority of research done up front
- More accurate time estimates
- More predictable release date

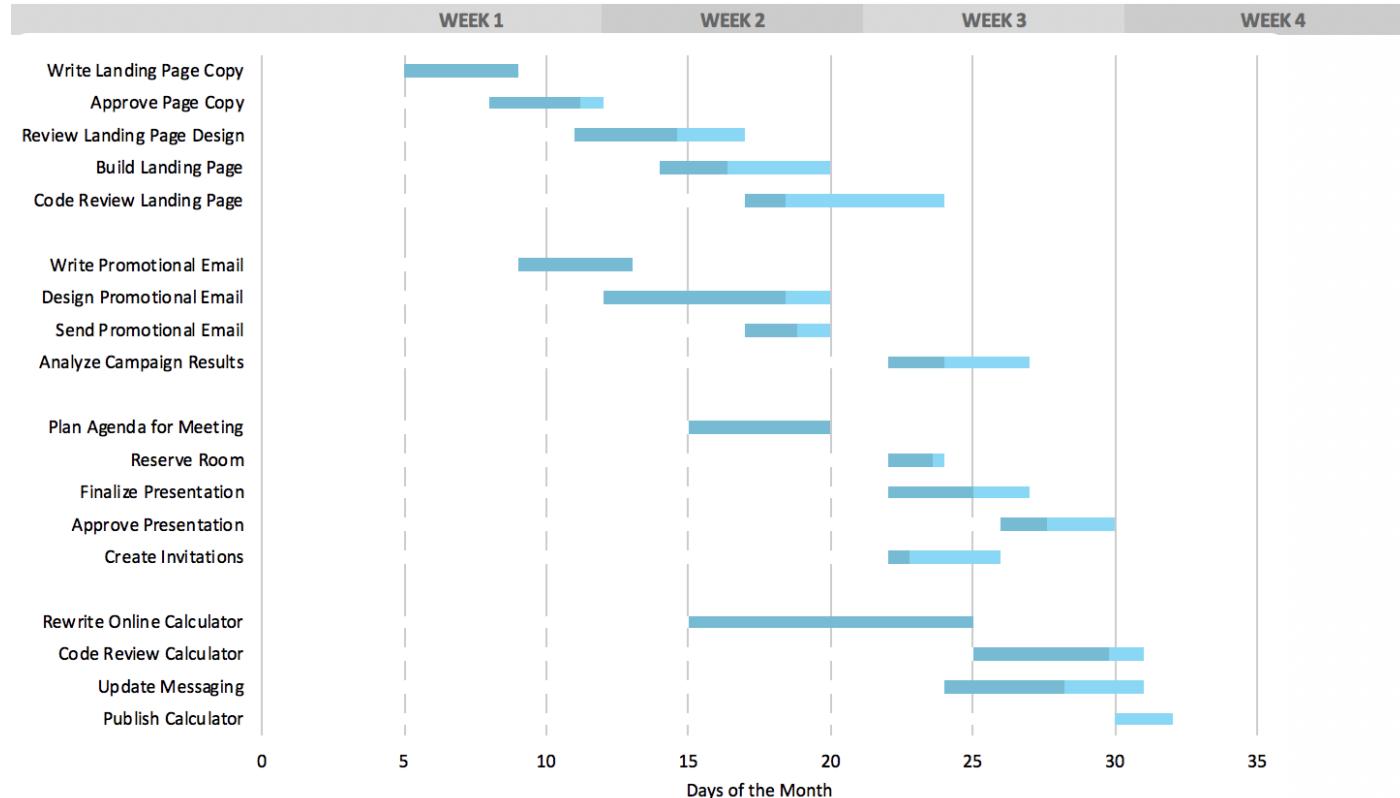
Cons:

- Process is brittle – can't pivot easily in terms of changing requirements
- Long lead times – difficult to respond to rapid business evolution

Tools:

- Commonly use Gantt charts to track projects, subtasks and dependencies

# Typical Gantt Chart



# Agile Software Development

A group of software development methodologies based on iterative development.

Requirements and solutions evolve through collaboration between self-organizing cross-functional teams.

Key aspects:

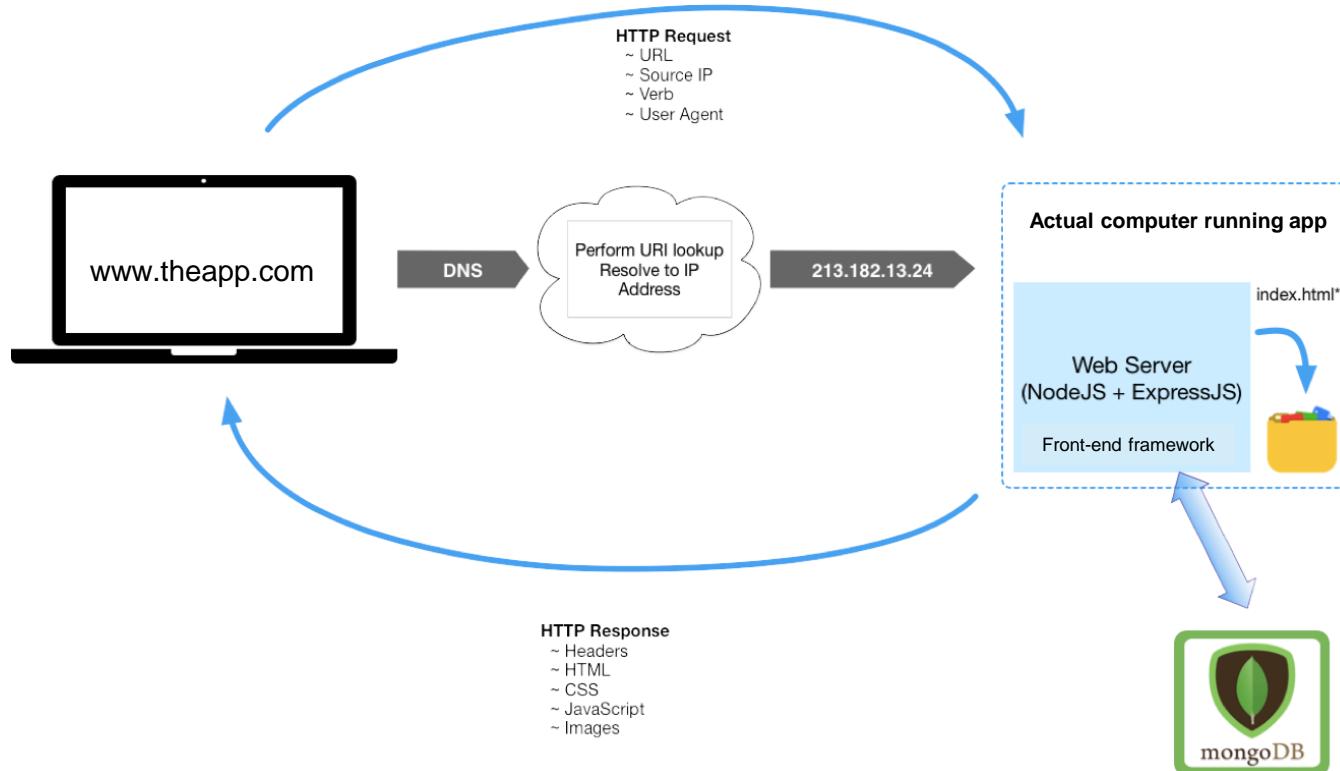
- Incremental delivery
- Always ready to ship
- Continuous inspection of work product and process provides feedback for continuous improvement

# Iterative Software Development

Key elements:

- Incremental improvement to code and process
- Always ready to ship
- Iterative feature development
- Working with the Product Owner
- Pivoting mid-development

# Elements of a Full-Stack App and HTTP Request/Response Cycle



## Group Discussion:

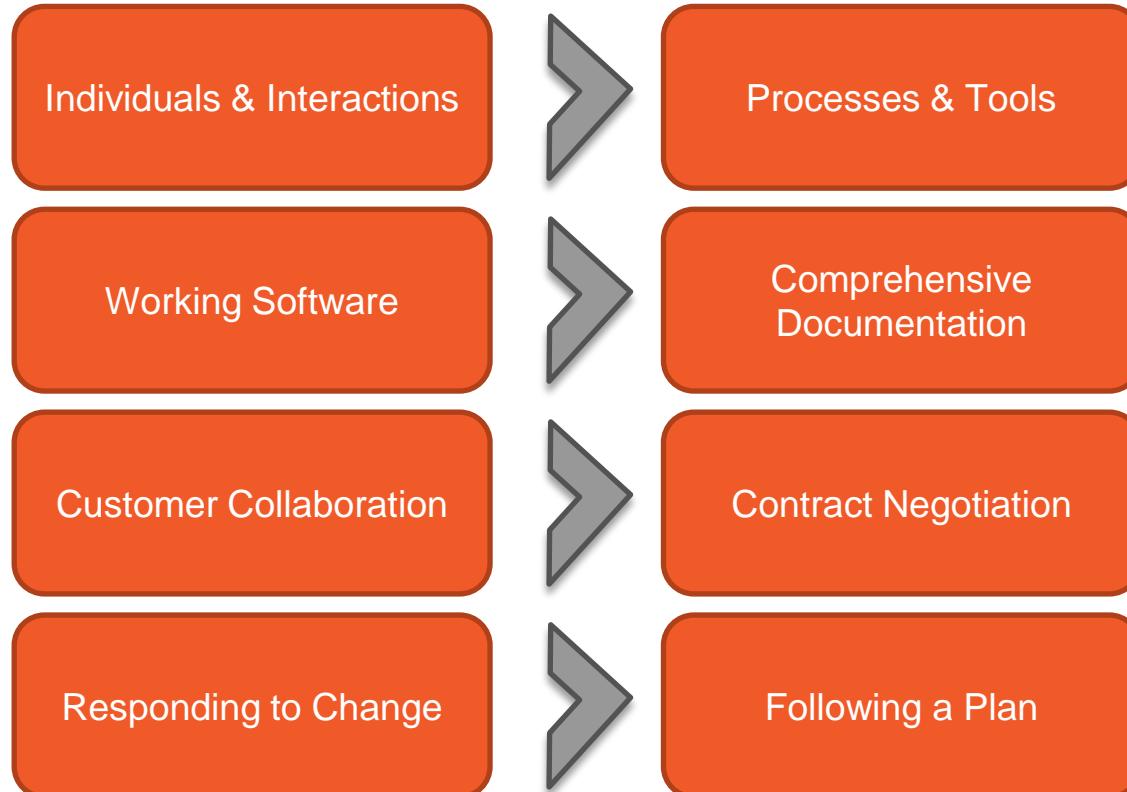
Looking at the  
“Reference  
Implementation” from  
an iterative approach

Group discussion considering the Reference Implementation from the following points of view:

- Itemize the set of high-level features that will be needed to implement this application?
- Which features should be implemented first? Which features should ship first?
- How might your approach need to pivot, based on market or other factors?

# Agile as a Development Methodology

# Agile Values



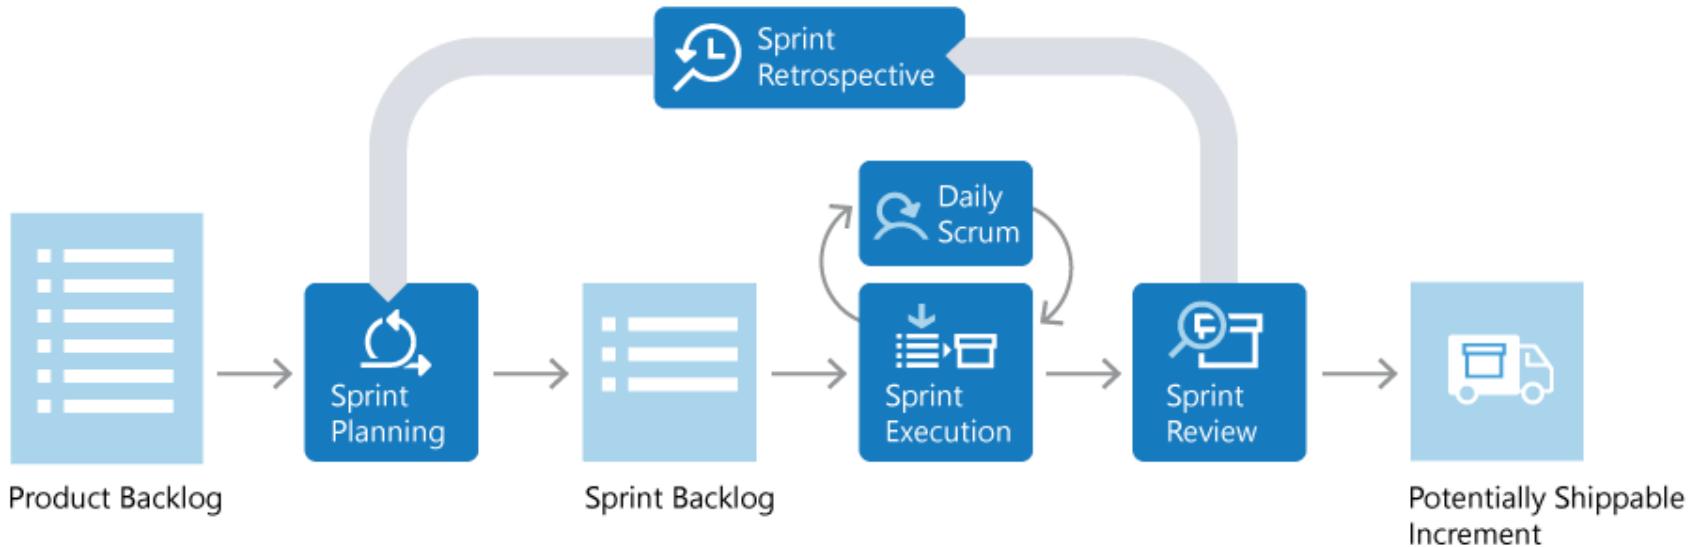
# Modern Approaches to Agile Software Development

Several project management approaches implement Agile values in some way. Often, they bring in elements of other approaches.

Popular methodologies:

Kanban	Scrum	Scaled Agile Framework (SAFe)
<ul style="list-style-type: none"><li>• Uses visual boards to view &amp; organize tasks</li><li>• Uses elements of “just in time” lean manufacturing strategies</li><li>• Emphasizes throughput</li></ul>	<ul style="list-style-type: none"><li>• Aligns closely with Agile values</li><li>• Breaks down product development into iterative sprints</li><li>• Makes use of exclusive roles (“Scrum Lead”, “Product Owner”)</li><li>• Emphasizes constant communication</li></ul>	<ul style="list-style-type: none"><li>• Workflow and organizational patterns to help deploy Agile at scale</li><li>• Can support large organizations</li></ul>

# Diagram of Typical Scrum Sprint



Source: <https://docs.microsoft.com/en-us/devops/plan/what-is-scrum>

# Scrum Roles

Product Owner

Scrum Lead

Scrum Team

# Scrum Roles

Product Owner

- Responsible for what team builds and why
- Keeps backlog up to date and in correct priority order

Scrum Lead

Scrum Team

# Scrum Roles

- Ensures that Scrum process is followed by team
- Responsible for the fidelity of the process
- Part coach, part team member, part cheerleader

Product Owner

Scrum Lead

Scrum Team

# Scrum Roles

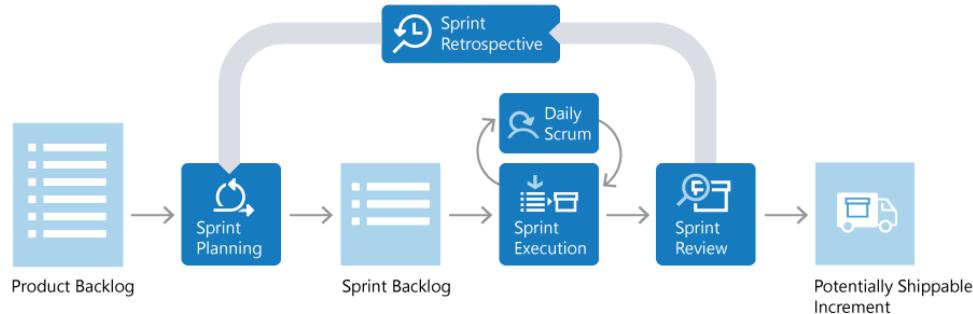
Product Owner

Scrum Lead

Scrum Team

- People building the product
- Responsible for product build (and associated quality)

# Group Reflection



- Where does architecture fit within each “ceremony”?
- How do we ensure architecture is not an afterthought but an active consideration during sprint planning & execution?

# Architecting for the Future

# SOLID Principles

SOLID principles help us build testable code

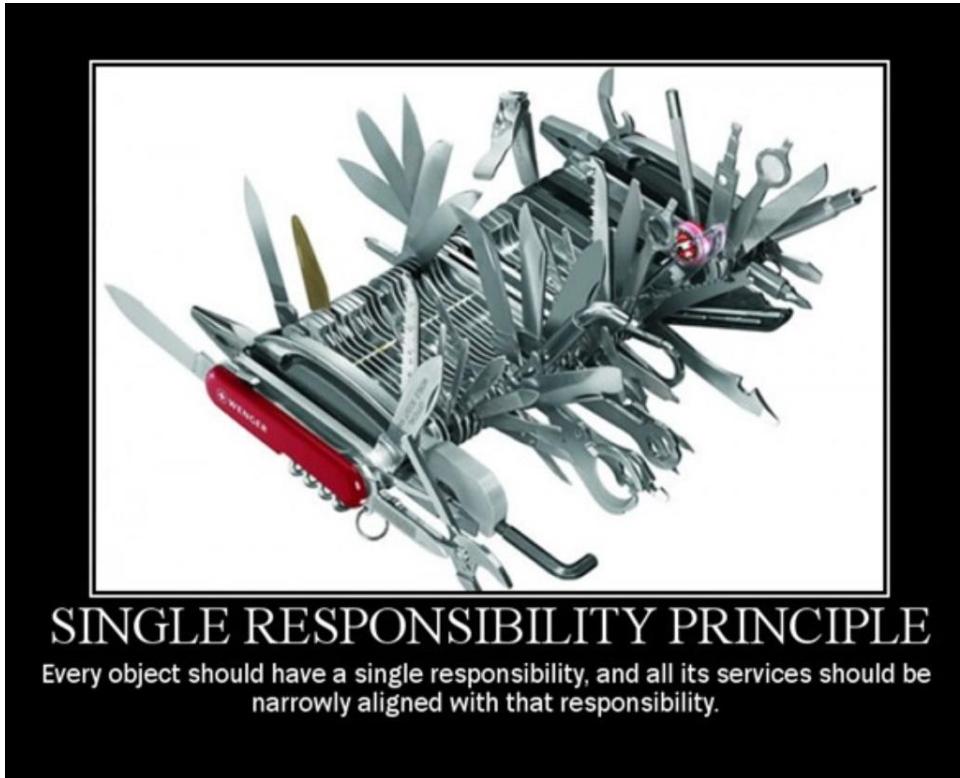
- Single Responsibility Principle (SRP)
- Open-Closed Principle (OCP)
- Liskov Substitution Principle (LSP)
- Interface Segregation Principle (ISP)
- Dependency Inversion Principle (DIP)

## Single Responsibility Principle (SRP)

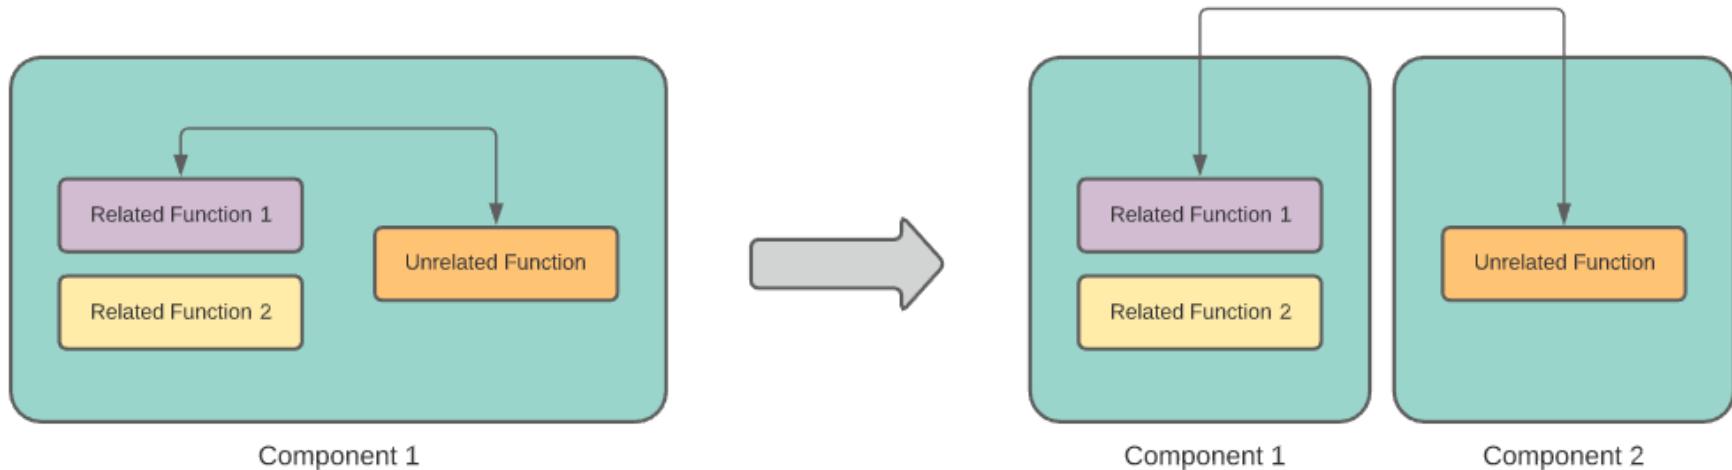
# Single Responsibility Principle (SRP)

*A system module or component should have only one reason to change*

# Single Responsibility Principle (SRP)



# Single Responsibility Principle (SRP)



## **LAB 01:**

Single Responsibility  
Principle (SRP)

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab01>

## Open-Closed Principle (OCP)

# Open-Closed Principle (OCP)

*Software entities should be open for extension but closed for modification*

## Open-Closed Principle (OCP)



OPEN CLOSED PRINCIPLE

Brain surgery is not necessary when putting on a hat.

## **LAB 02:**

Open Closed Principle  
(OCP)

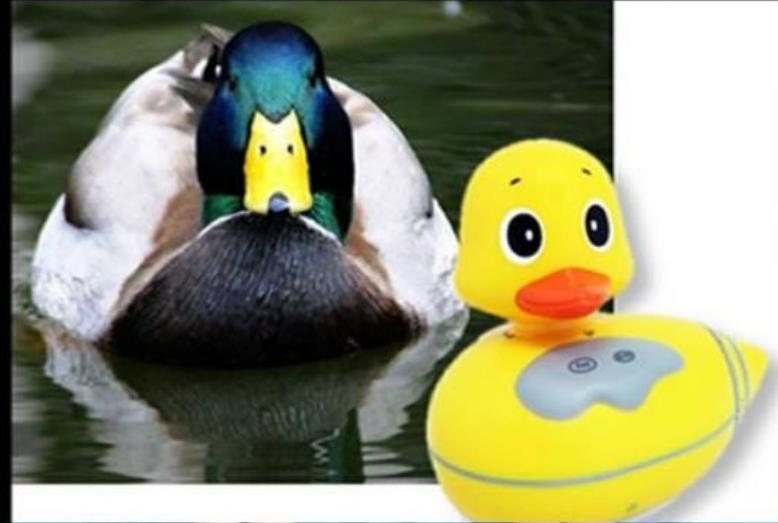
<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab02>

## Liskov Substitution Principle (LSP)

# Liskov Substitution Principle (LSP)

*Subtypes must be substitutable for their base types*

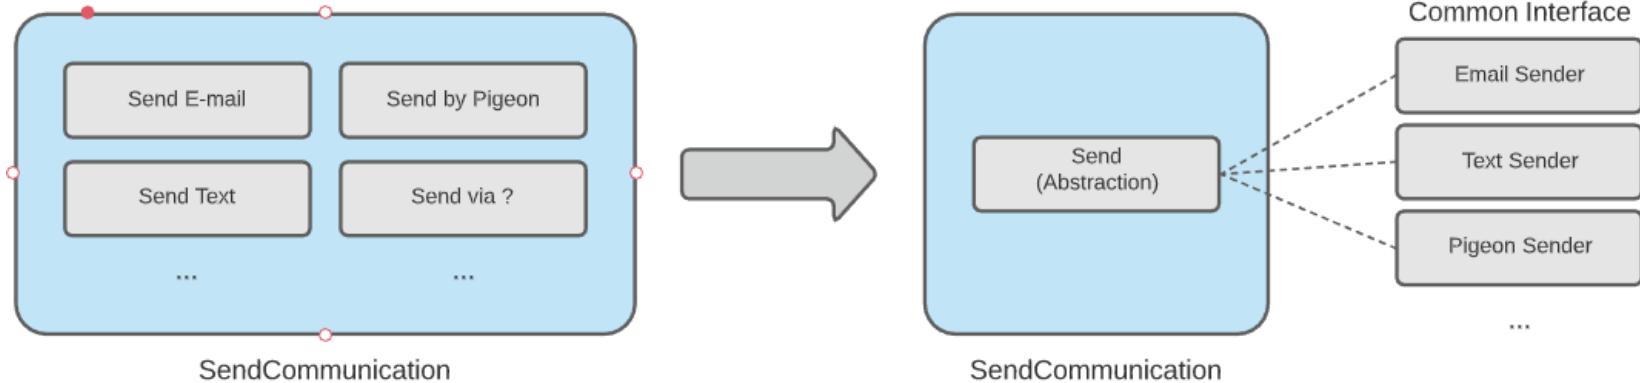
## Liskov Substitution Principle (LSP)



### Liskov Substitution Principle

If it looks like a duck and quacks like a duck but needs batteries,  
you probably have the wrong abstraction.

# OCP & LSP



## **LAB 03:**

Liskov Substitution  
Principle (LSP)

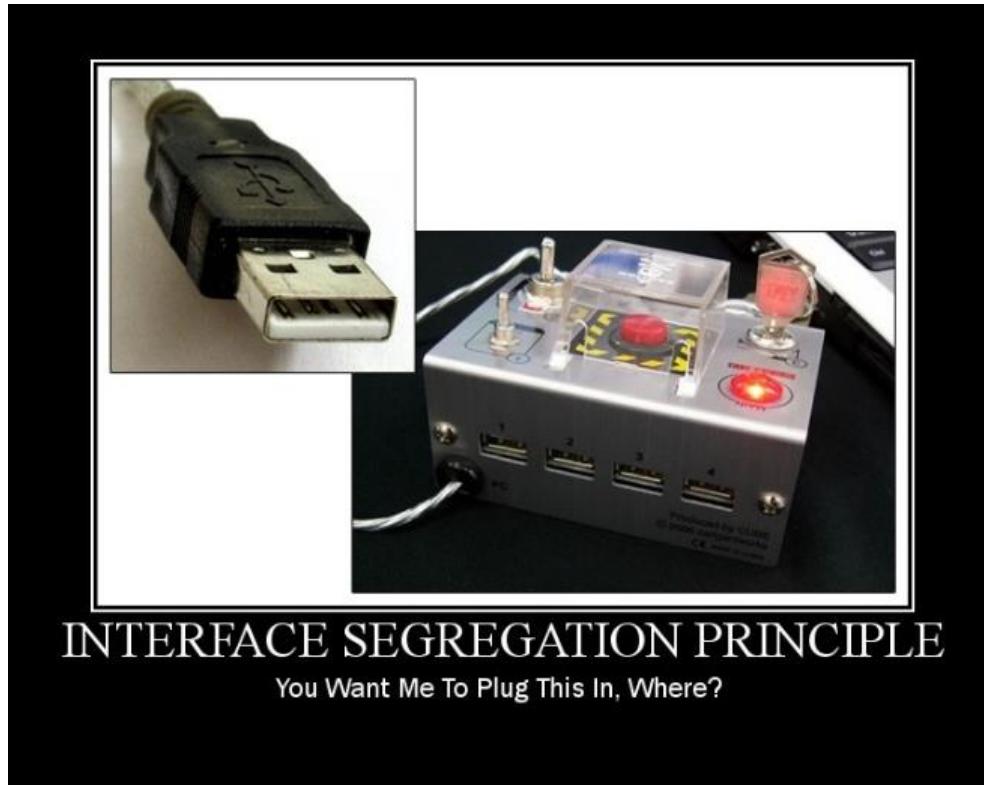
<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab03>

## Interface Segregation Principle (ISP)

# Interface Segregation Principle (ISP)

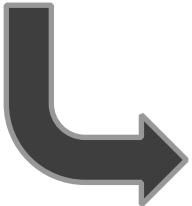
*Clients should not be forced to depend on methods they do not use*

# Interface Segregation Principle (ISP)



# Interface Segregation Principle (ISP)

```
6  
7     namespace CoolCompany.Financials  
8     {  
9         public interface ITaxProcessor  
10        {  
11            double CalculateSalesTax(double onAmount);  
12            double CalculatePropertyTax(double onAmount);  
13            double CalculateIncomeTax(double onAmount);  
14        }  
15    }
```



```
6  
7     namespace CoolCompany.Financials  
8     {  
9         public interface ITaxProcessor  
10        {  
11            double Calculate(double onAmount);  
12        }  
13    }  
14    public interface ISalesTaxProcessor : ITaxProcessor  
15    {  
16        double LookupStateTaxRate(string state);  
17    }  
18    public interface IPropertyTaxProcessor : ITaxProcessor  
19    {  
20        double LookupTownshipTaxRate(string township);  
21    }  
22    public interface IIIncomeTaxProcessor : ITaxProcessor  
23    {  
24        double CalculateBackTaxes(string taxpayerId);  
25    }  
26    }  
27 }  
28 }
```

## **LAB 04:**

Interface Segregation  
Principle (ISP)

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab04>

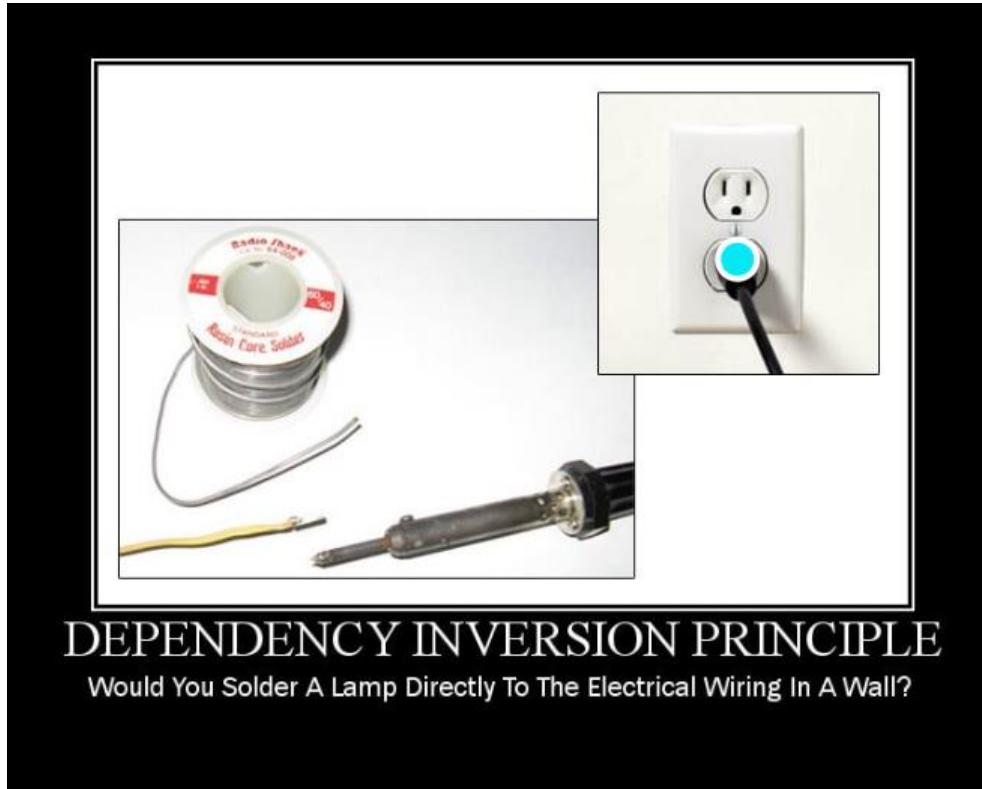
## Dependency Inversion Principle (DIP)

# Dependency Inversion Principle (DIP)

*High-level modules should not depend on low-level modules – both should depend on abstractions*

*Abstractions should not depend upon details – details should depend upon abstractions*

# Dependency Inversion Principle (DIP)



## **LAB 05:**

### Dependency Inversion Principle (DIP)

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab05>

# Test-Driven Development (TDD)

# Benefits to be experienced with automated testing



Greater shared success  
(as tests can be tied to  
requirements)



Cleaner code



Guaranteed system  
documentation



Improved quality



Reduced fear



Improved regression  
protection

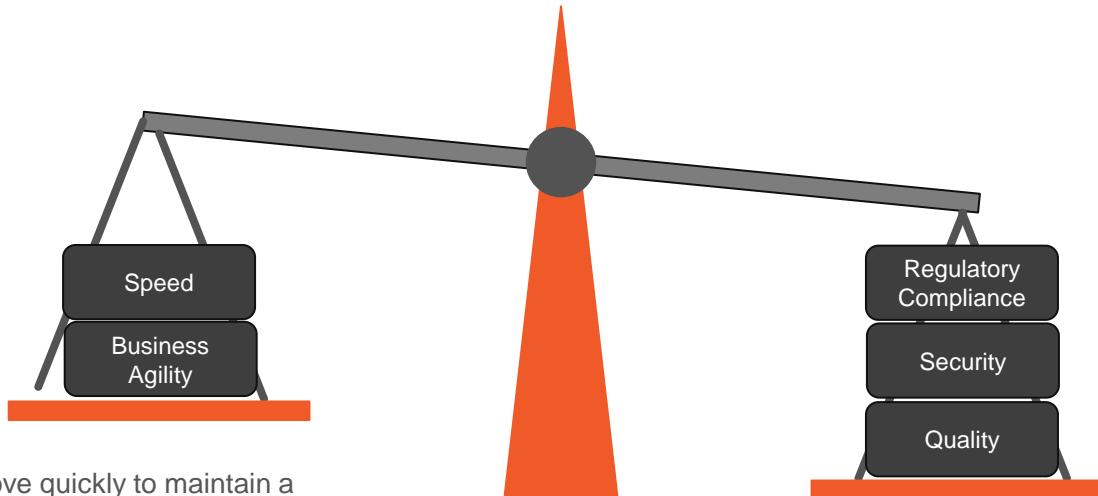


Greater streamlining to the  
approach  
(small bites, avoidance of  
premature optimization,  
structured flow)



Discussion | Which of these benefits resonate most with you?

# Balancing Speed & Quality



We want to move quickly to maintain a competitive advantage & increase profitability

We want to move safely to ensure we minimize rework cost & prevent monetary & reputational damage

# Test-Driven Development (TDD) and the 3 Laws

1

2

3

## *The First Law*

“You are not allowed to write any production code until you have written a unit test that fails due to its absence”

## *The Second Law*

“You are not allowed to write more of a unit test than is sufficient to fail, and failing to compile is failing”

## *The Third Law*

“You are not allowed to write more production code than is sufficient to cause the currently failing test to pass”

# Understanding TDD

Scenario: Determine discount based on purchase volume



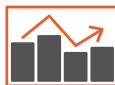
- SaaS module is external service defining and enforcing business rules
- Several levels of discount based on purchase volume ranges
- Need to verify that the invoicing module operates correctly with the different levels of discount
- What if the call to the SaaS module fails?



- Tie test cases directly to specific business requirements (e.g., each discount level)
- Write **just enough** code to pass the tests
- By extension, have written **just enough** code to meet the business requirements
- **Mocking** enables tight control over what gets returned from the downstream dependency, including any error conditions
- Result is verified quality when built and **automated regression testing** against future changes

# Best practices with TDD

TDD supports the following best practices:



Don't commit to a methodology, commit to impact and a positive result



Ensure that the engineering team has a clear understanding (and practice) for unit vs. integration testing



Automated test code is code – just like application code – estimate and plan accordingly



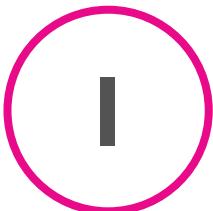
Find the right % of test coverage – quality & quantity!

# To be effective, automated unit tests need to follow FIRST:



Fast

or they won't be  
executed



Independent

i.e., minimal  
dependencies  
between tests (keeps  
them simpler)



Repeatable

i.e., every run against  
unchanged code  
should operate the  
same



Self-validating

i.e., Boolean output  
(pass or fail)



Timely

if not first, early and  
often

## Challenges

What are some of the challenges you face with software testing in your jobs (i.e., why don't we test as much as we want to or should)?

## Challenges

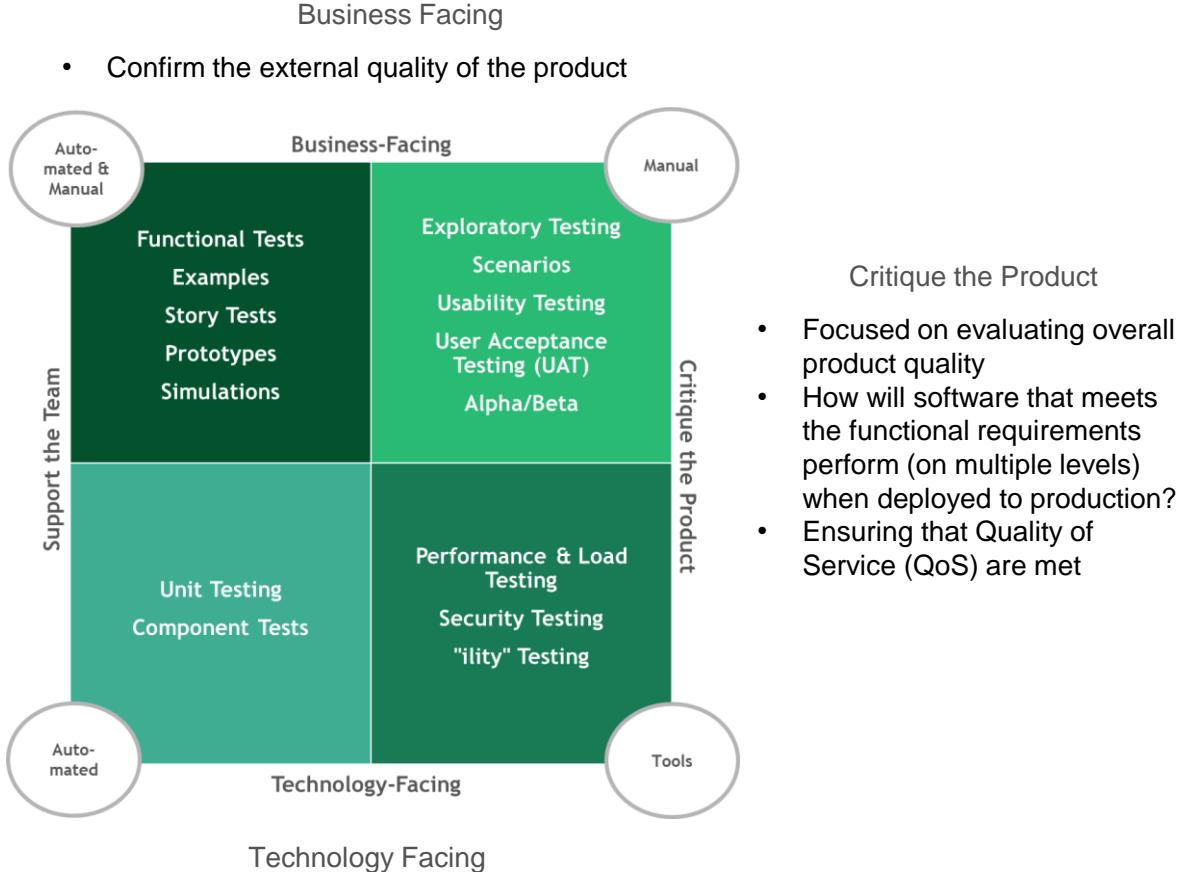
- Time pressures
- Complexity
- Requires more code
- More debugging (up front)
- More cost (up front)
- “Can’t we just ensure quality with other types of testing?”

# Testing Strategies

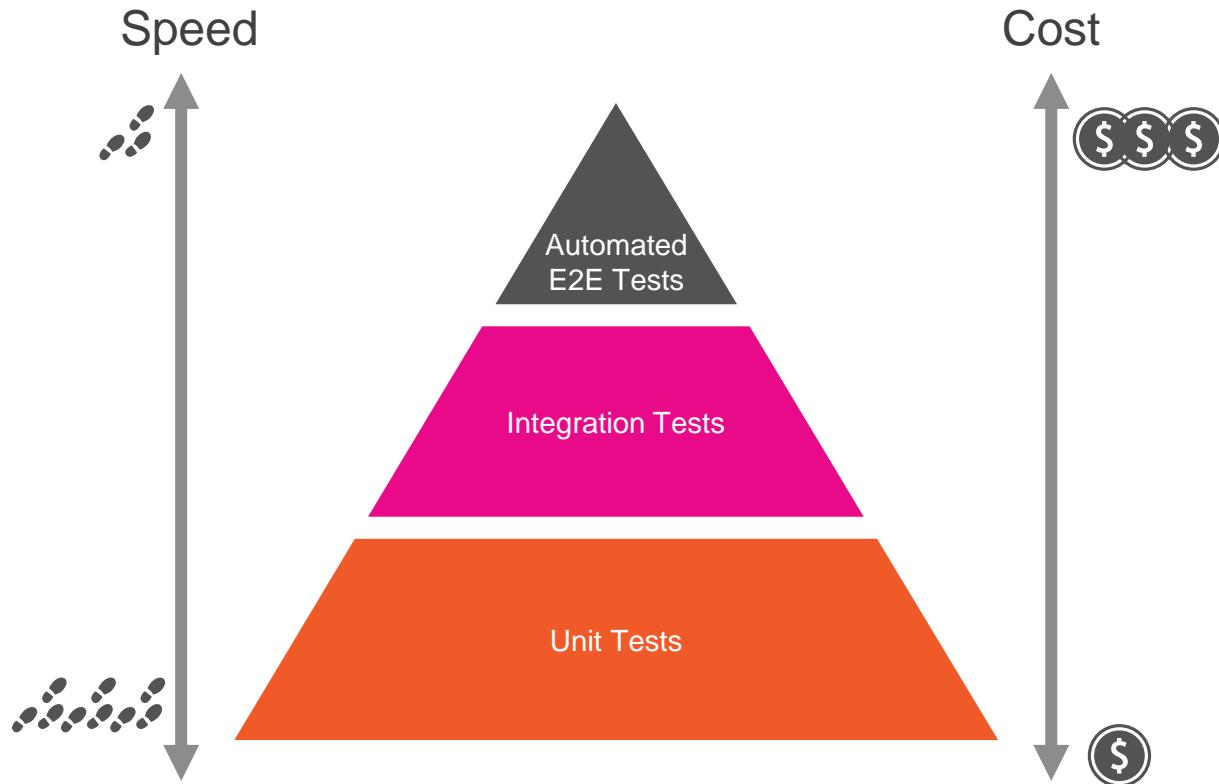
# Testing Quadrant

## Support the Team

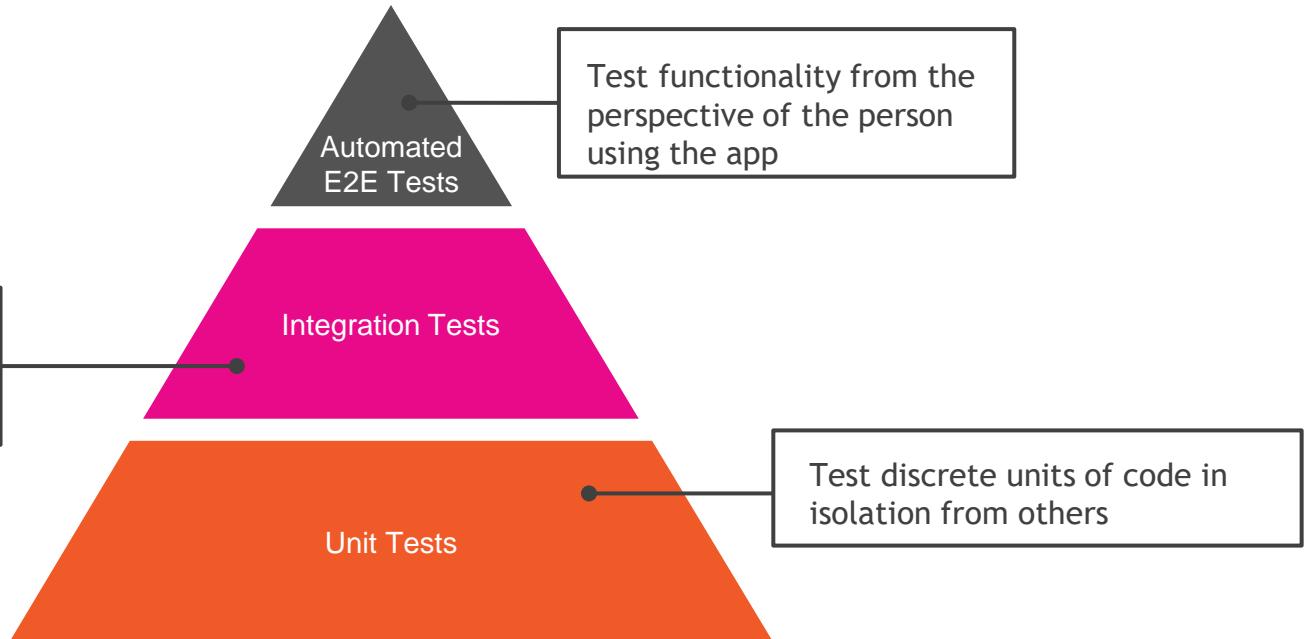
- Focused on supporting the teams that underpin the product & its delivery
- Goal is to help guide the quality of the product & prevent functional defects from making their way into production



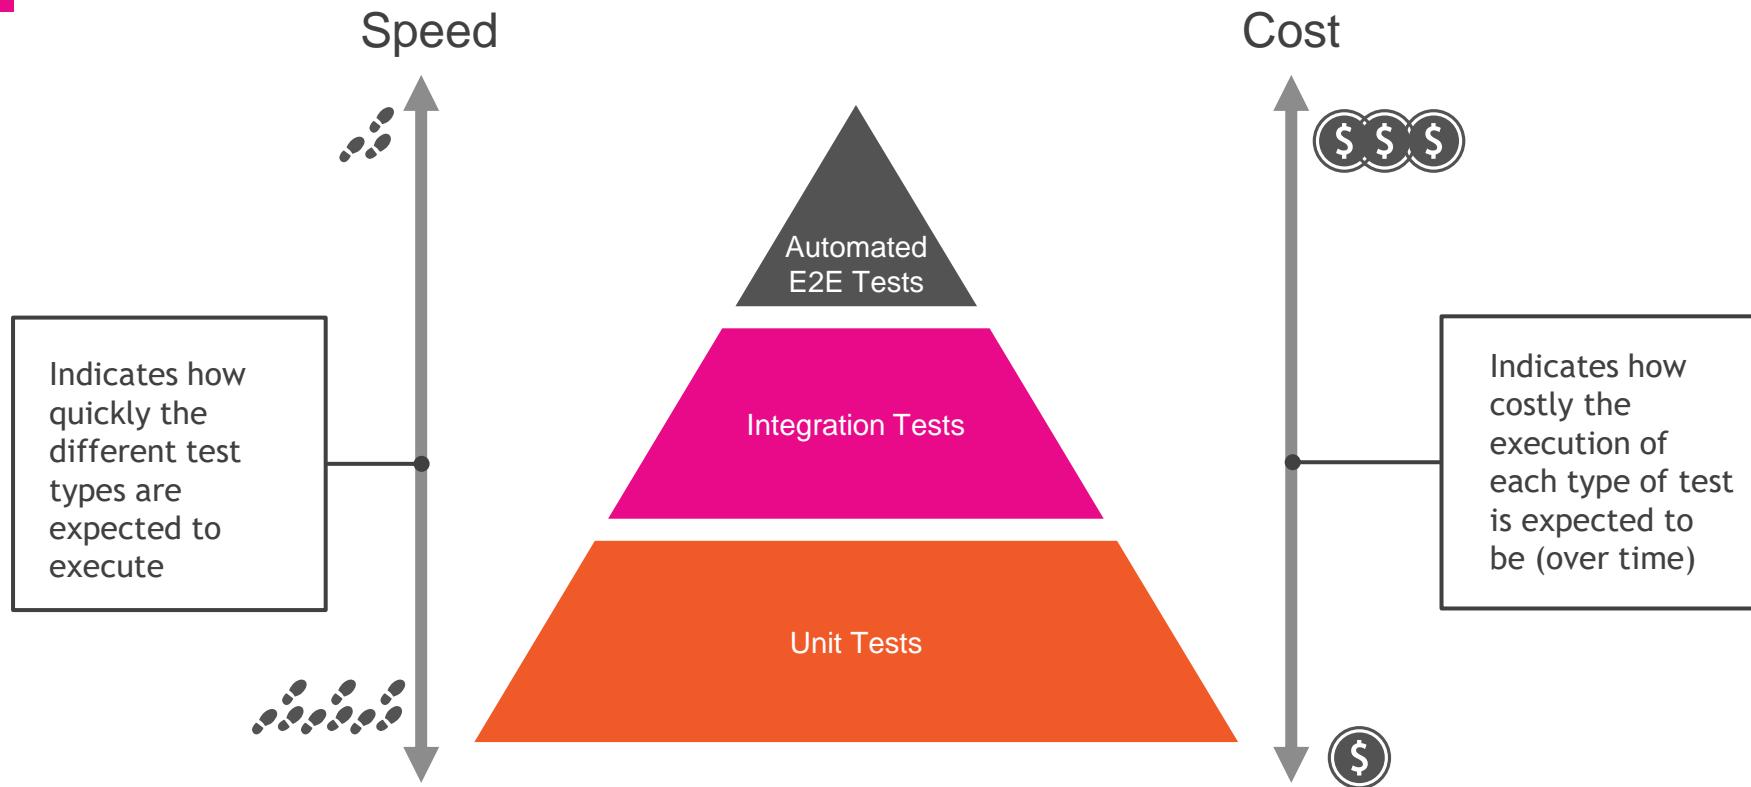
# Test Pyramid



# Test Pyramid



# Test Pyramid



## **LAB 06:**

### Testing Classes

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab06>

## **LAB 07:**

### Testing Classes

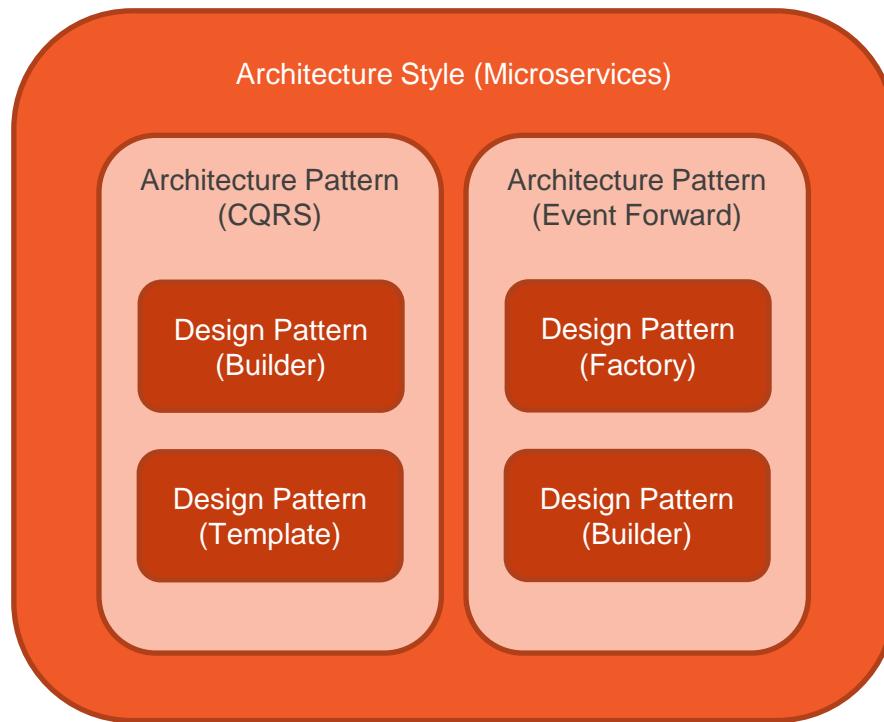
<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab07>

## LAB 08: TDD

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab08>

# Software Architecture Styles

# Redux



# Common Software Architecture Styles

Layered Architecture

Microkernel Architecture

Event-Driven  
Architecture

Microservices  
Architecture

# Technical vs. Domain Partitioning



## Technical

- Components organized by technical usage
- Elements of a given business domain are spread across multiple components or layers
- Useful if change isolated to a specific technical area
- Includes Layered, Microkernel, Event-Driven, and Space-Based

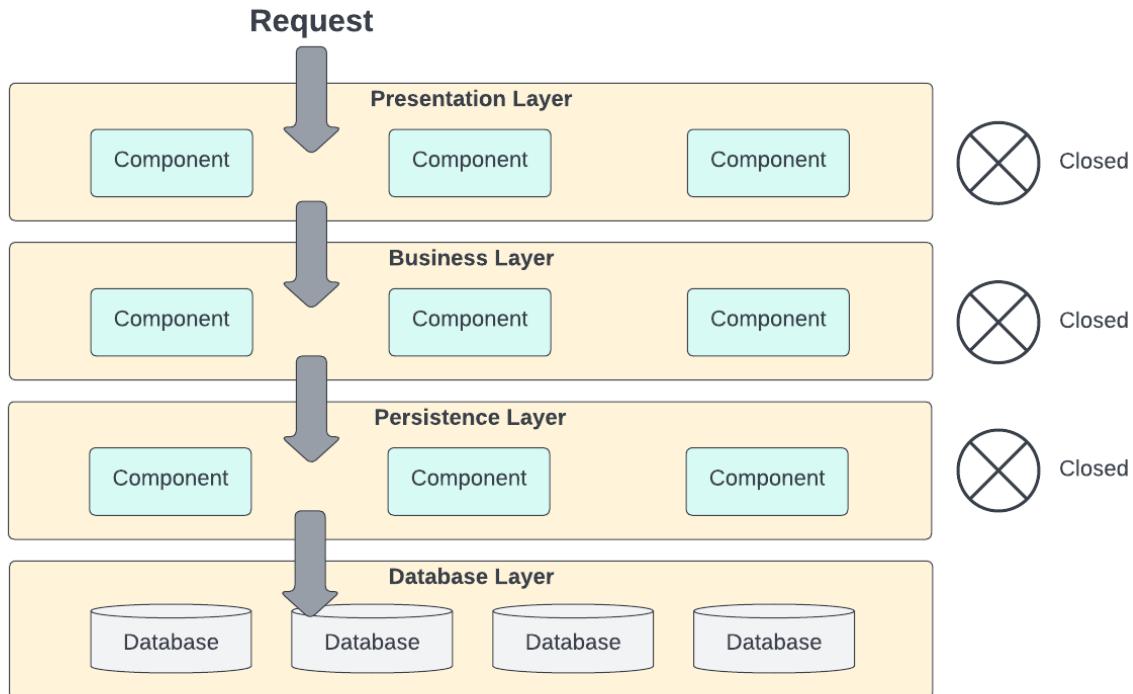


## Domain

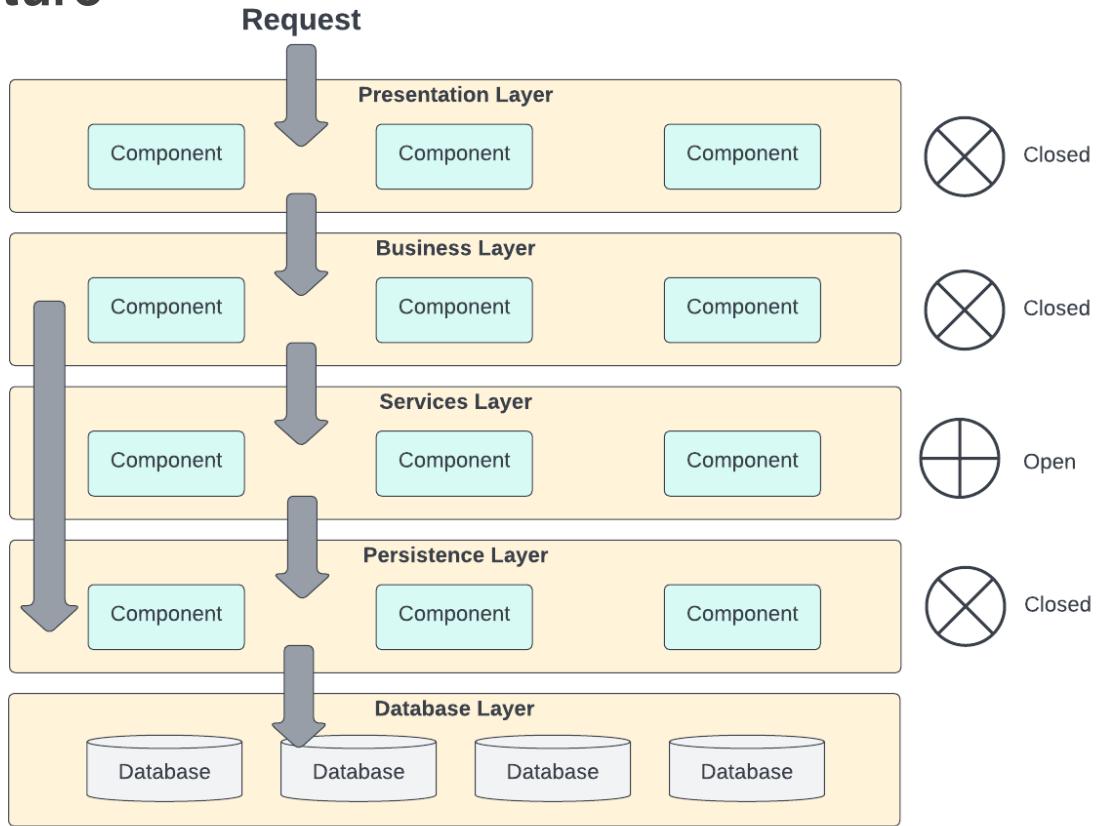
- Organized by business domain area vs. technical usage
- AKA Domain-Driven Design (DDD)
- Changes typically isolated to an individual domain
- Helps align technical implementation with business processes
- Includes Microkernel and Microservices

# Layered Architecture

# Layered Architecture



# Layered Architecture



## Group Discussion:

### Layered Architecture

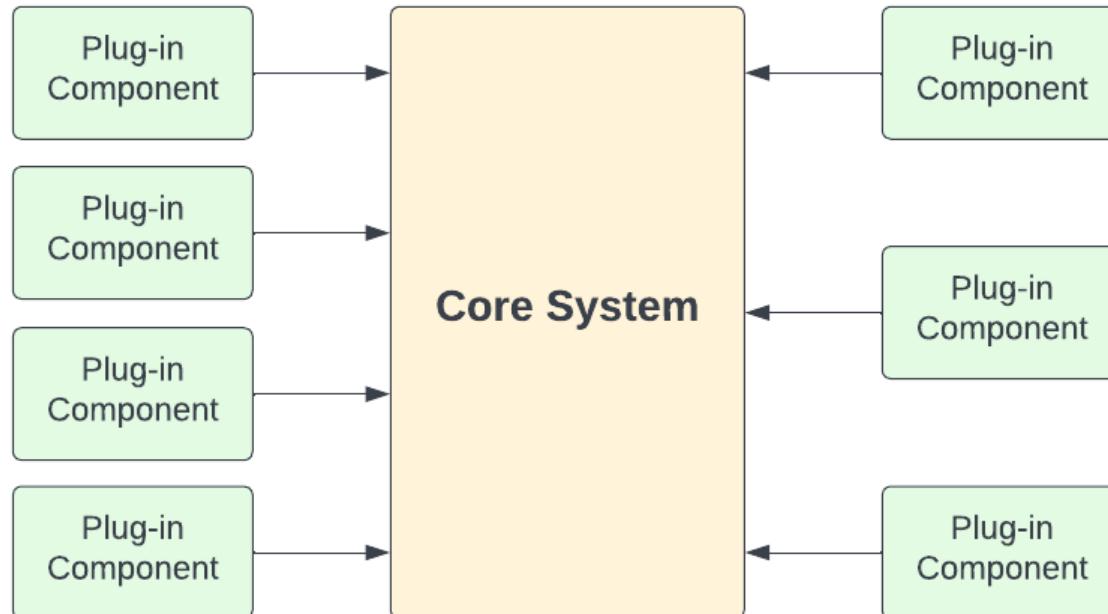
As a group, discuss the pros & cons of this architectural style.  
Specifically:

- What are some advantages to using this style? Is there a specific application profile that would cause you to prefer this style?
- What are some disadvantages to using this style? Is there a specific application profile that would you to avoid this style?

# Microkernel Architecture

# Microkernel Architecture

Independent, standalone modules containing specialized logic to address one or more specific use cases or provide specialized business capability



Minimal functionality required to make base system operational

# Microkernel Architecture

## Key Components & Strategies of Approach

-  Keep cross-plugin communication to a minimum
-  Core system uses a separate plug-in registry to know what's registered and available
-  Implemented as separate libraries or modules (can also be implemented remotely)

## Group Discussion:

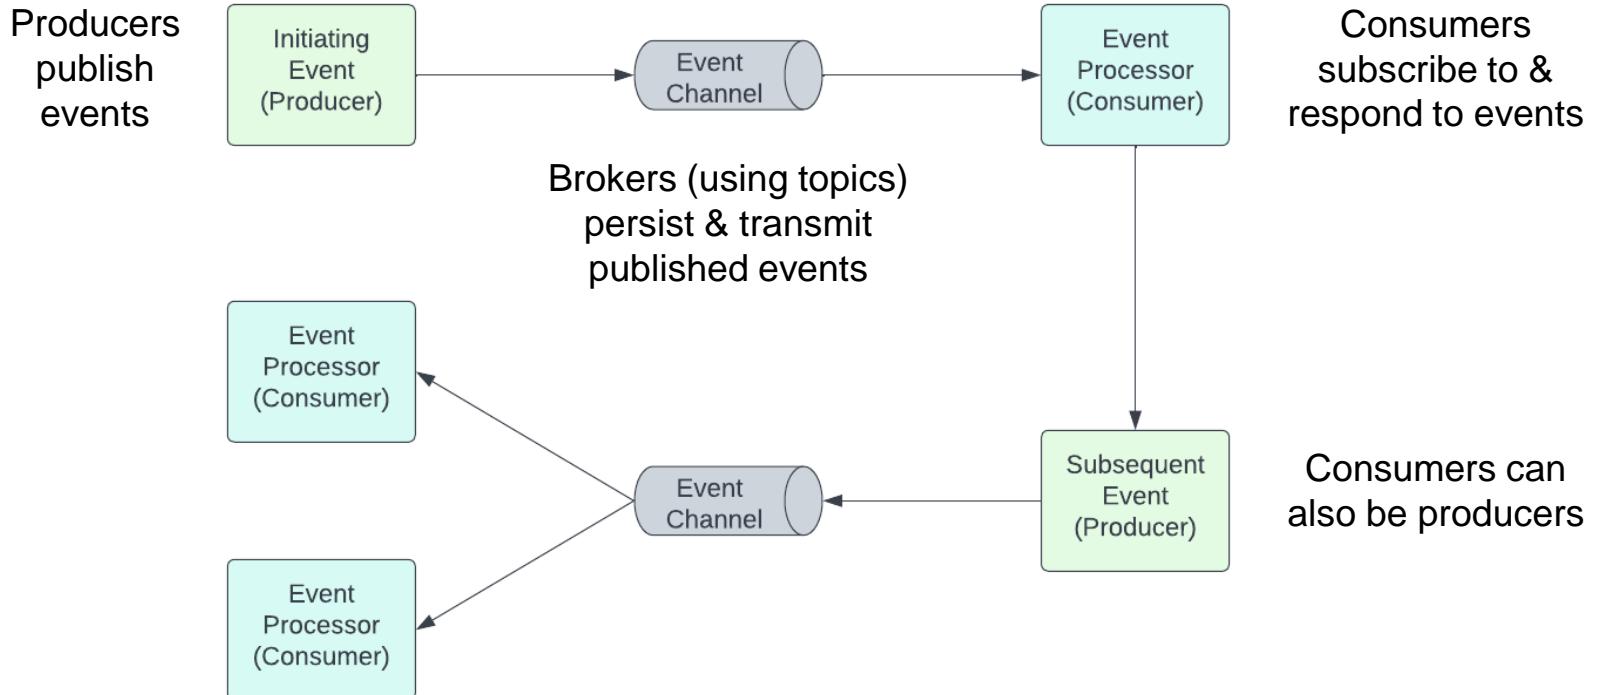
Microkernel Architecture

As a group, discuss the pros & cons of this architectural style.  
Specifically:

- What are some advantages to using this style? Is there a specific application profile that would cause you to prefer this style?
- What are some disadvantages to using this style? Is there a specific application profile that would cause you to avoid this style?

# Event-Driven Architecture

# Event-Driven Architecture





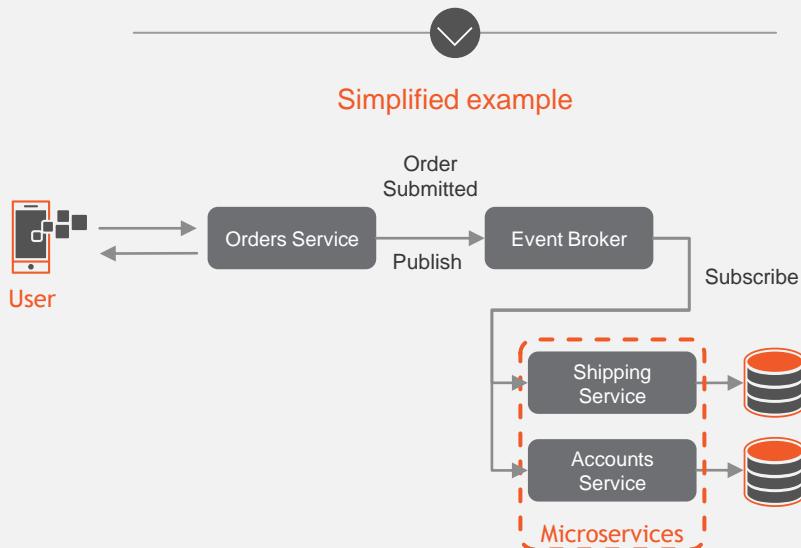
# Event-Driven Architectures



As an architectural style:

- Enables looser coupling between communicating services
- Provides dynamic flexibility in the definition of integrated workflows
- Promotes resiliency as an intermediate component (the broker) is used to separately track and manage the queue events
- Uses concept of domain events – publish of events with business significance

## Simplified example



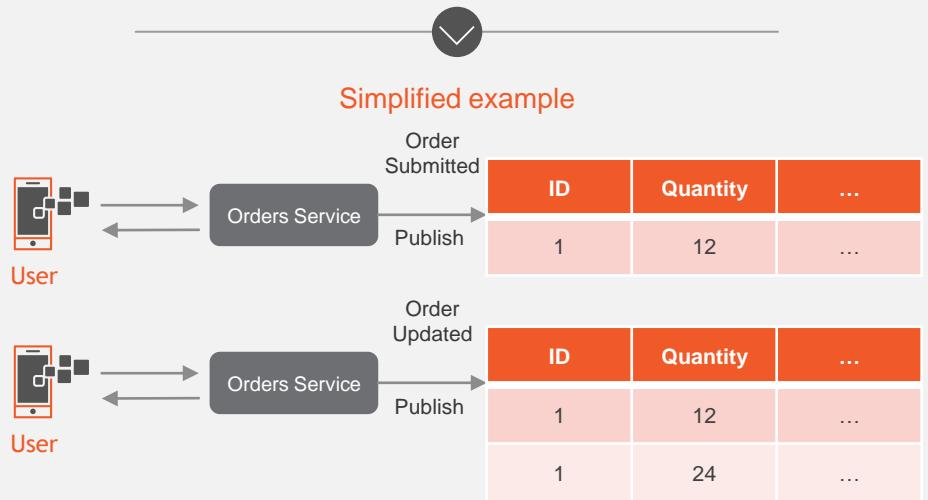


# Event-Driven Architectures



## Event Sourcing

- Data changes are captured as immutable events
- When a change is made, rather than overwrite existing record, a new record is created to reflect updated state
- Provides a full picture of what happened during an entity's lifecycle – no data destruction
- Can make reads more difficult



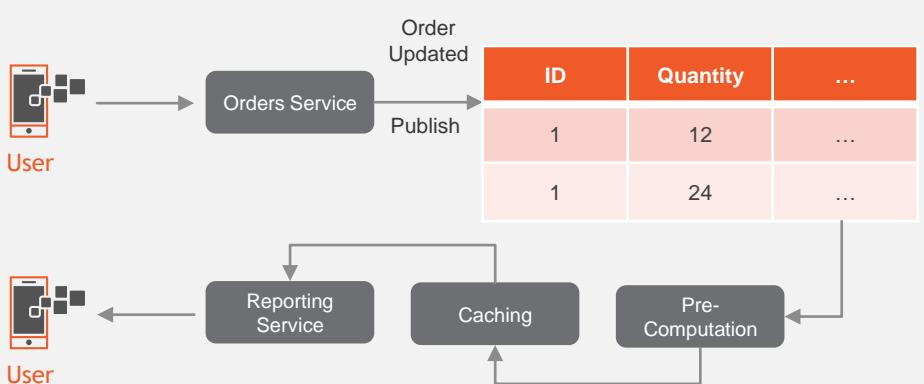


# Event-Driven Architectures



## CQRS

- Command-Query Responsibility Segregation
- Can help address challenges with Event Sourcing architectures
- Separate processes used for writes and reads
- Allows write optimization (“raw” events) and read optimization (dedicated aggregation)
- Also, enables denormalization and reformatting for reads separate from write structures



# Domain Events vs. Commands (or Messages)

## Domain Events

- Notification that something of business significance happened
- Can range from 0 or more processes/components (consumers) that “care”

## Commands

- Explicit message intended for a specific target component
- Used to initiate a new step in a flow
- Can be combined with domain events

## Group Discussion:

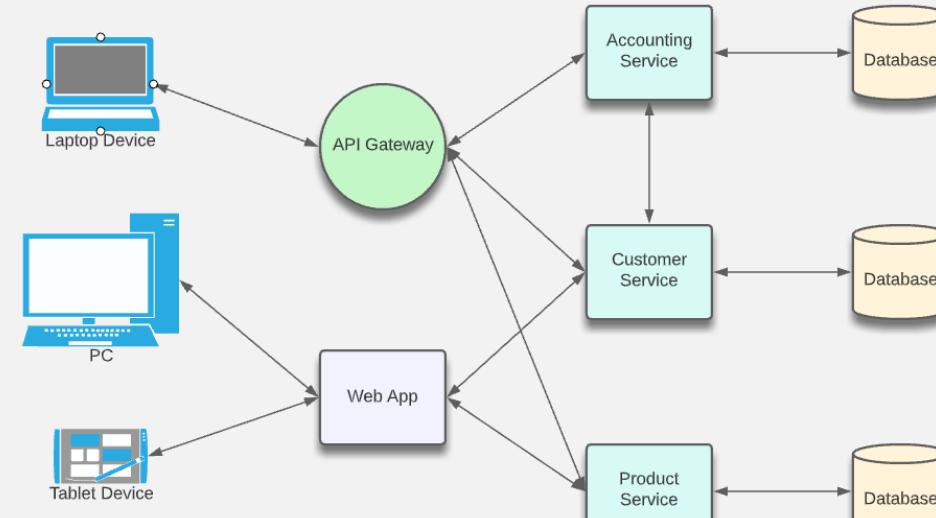
### Event-Driven Architecture

As a group, discuss the pros & cons of this architectural style.  
Specifically:

- What are some advantages to using this style? Is there a specific application profile that would cause you to prefer this style?
- What are some disadvantages to using this style? Is there a specific application profile that would cause you to avoid this style?

# Microservices Architecture

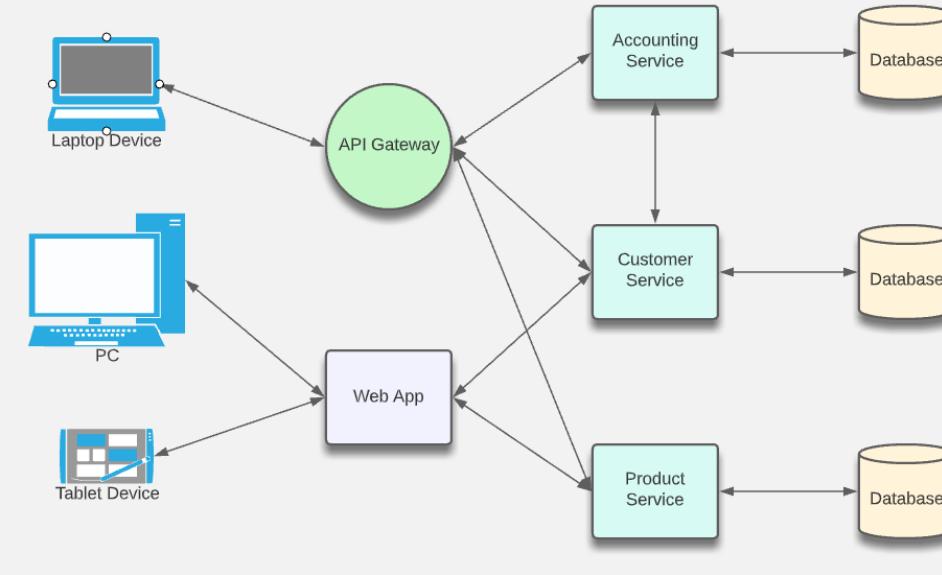
# Microservices Architecture



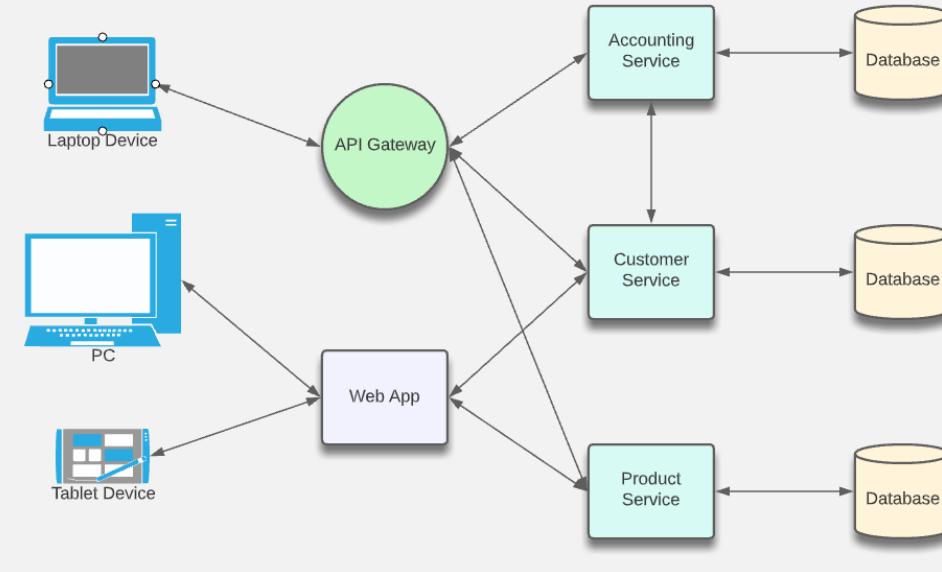
# Microservices Architecture

## Benefits

- Enables work in parallel
- Promotes organization according to business domain
- Advantages from isolation
- Flexible in terms of deployment and scale
- Flexible in terms of technology
- Allows multiple agile teams to maximize autonomy in effort and technology



# Microservices Architecture



## Implications

- Requires a different way of thinking
- Complexity moves to the integration layer
- Organization needs to be able to support re-org according to business domain (instead of technology domain)
- With an increased reliance on the network, you may encounter latency and failures at the network layer
- Transactions must be handled differently (across service boundaries)

# Where Microservices & APIs Overlap



## APIs:

- Model the exchange of information between a consumer & producer
- Support business integration between two internal systems (“in process”) or between an internal system & an external partner (“out of process”)
- Built around technology specifications for interaction & type definition
- Enable integration over the network/Internet

# Technology Options for Building APIs



When deciding on a technology to use for building an API, today's landscape offers multiple options:

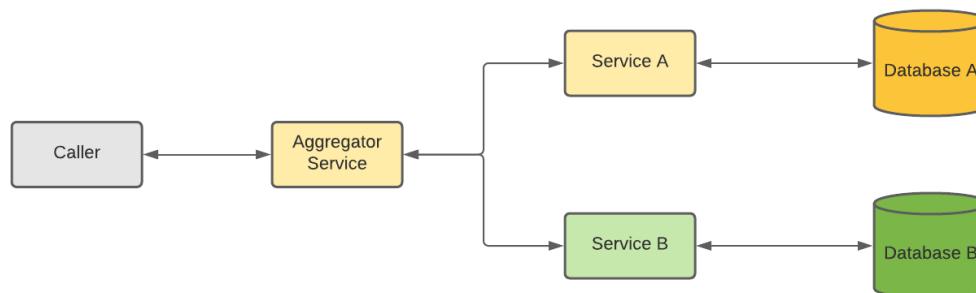
- ✓ Implementation using a serverless platform like AWS Lambda
- ✓ Exposing via Robotic Process Automation (RPA)
- ✓ Exposing as a set of endpoints using the Mulesoft platform (or similar)
- ✓ Implementation using a microservices-based approach & architecture

# Microservice Design Patterns

- Aggregator
- API Gateway
- Chain of Responsibility
- Asynchronous Messaging
- Circuit Breaker
- Anti-Corruption Layer
- Strangler Application
- Others as well  
(<https://microservices.io/patterns/index.html>)

# Microservice Design Patterns - Aggregator

- Akin to a web page invoking multiple microservices and displaying results of all on single page
- In the pattern, one microservice manages calls to others and aggregates results for return to caller
- Can include update as well as retrieval ops



# Microservice Design Patterns - Aggregator

## Benefits:

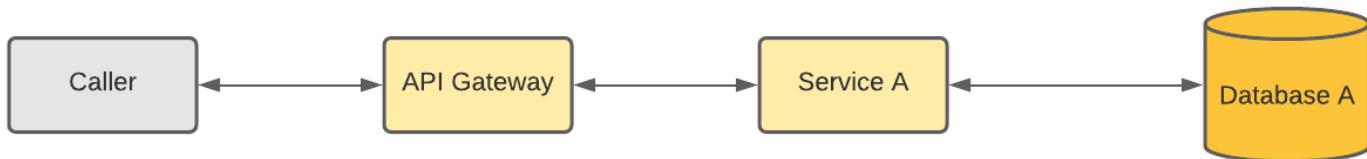
- Helps you practice DRY (Don't Repeat Yourself)
- Supports logic at time of aggregation for additional processing

## Potential “gotchas”:

- Performance if downstream microservices not called asynchronously

# Microservice Design Patterns – API Gateway

- Specific type of infrastructure that helps manage the boundary
- Provides an intermediary for routing calls to a downstream microservice
- Provides a protection and translation layer (if calling protocol different from microservice)
- Can be combined with other patterns as well
- Considered a best practice from a security perspective as well



# Microservice Design Patterns – API Gateway

## Benefits:

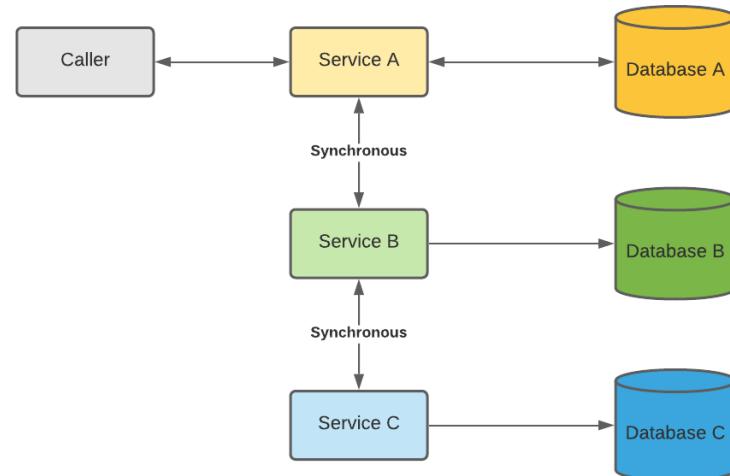
- Usually, built in throttling (to protect against DDoS)
- Can support translation between caller and downstream microservice for protocol, message format, etc.
- Provides a centralized point of entry for AuthN/AuthZ
- Can be combined with load balancing to enable scalability and resiliency
- Centralized logging

## Potential “gotchas”:

- Can increase cost
- Requires additional configuration & management but enables additional capabilities as well

# Microservice Design Patterns – Chain of Responsibility

- Represents a chained set of microservice calls to complete a workflow action
- Output of 1 microservice is input to the next
- Uses synchronous calls for routing through chain



# Microservice Design Patterns – Chain of Responsibility

## Benefits:

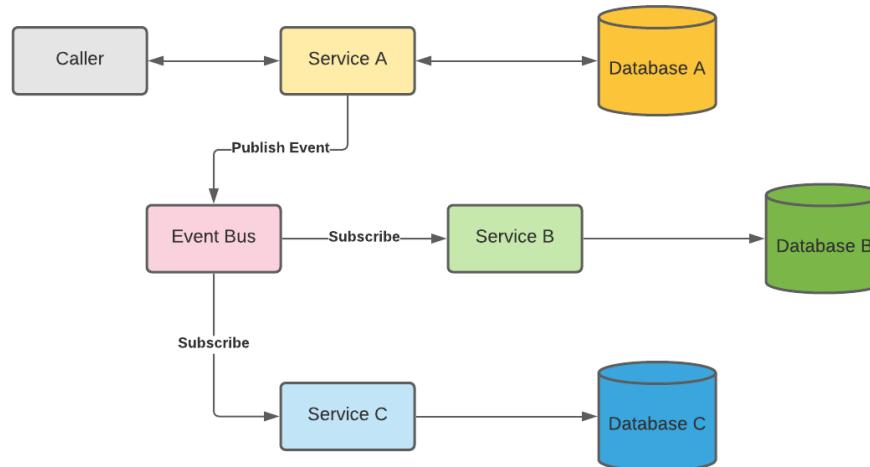
- Supports composition of microservices in a workflow that must happen in sequence
- Synchronous communication typically less complex

## Potential “gotchas”:

- Increased response time – response time becomes the sum of each microservice’s response time in the chain

# Microservice Design Patterns – Asynchronous Messaging

- Supports asynchronous interaction between independent microservices
- Messages published to a topic by a producer
- Topic subscribed to by one or more consumers



# Microservice Design Patterns – Asynchronous Messaging

## Benefits:

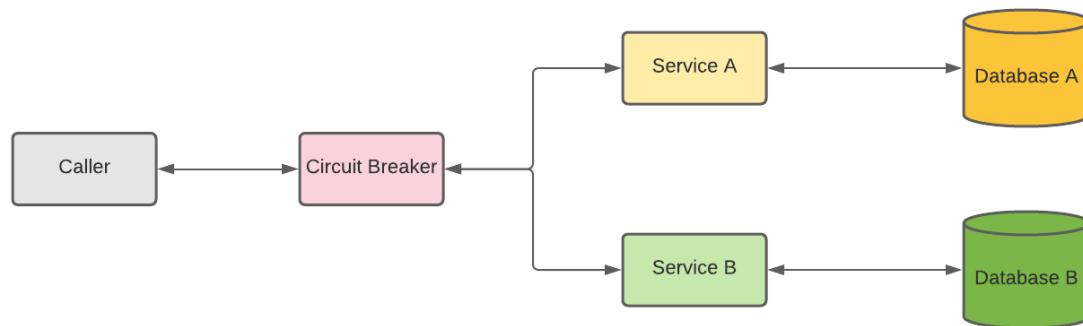
- Promotes looser coupling in cases where synchronous calls are not required
- Allows multiple services the option of being notified (vs. point-to-point)
- If given consumer is down, producer can continue to send messages and they won't be lost

## Potential “gotchas”:

- Additional complexity
- Can be harder to trace an action end-to-end (but there are ways to handle)
- Not a good fit if specific timing and sequencing required

# Microservice Design Patterns – Circuit Breaker

- Prevents unnecessary calls to microservices if down
- Circuit breaker monitors failures of downstream microservices
- If number of failures crosses a threshold, circuit breaker prevents any new calls temporarily
- After defined time period, sends a smaller set of requests and ramps back up if successful



# Microservice Design Patterns – Circuit Breaker

## Benefits:

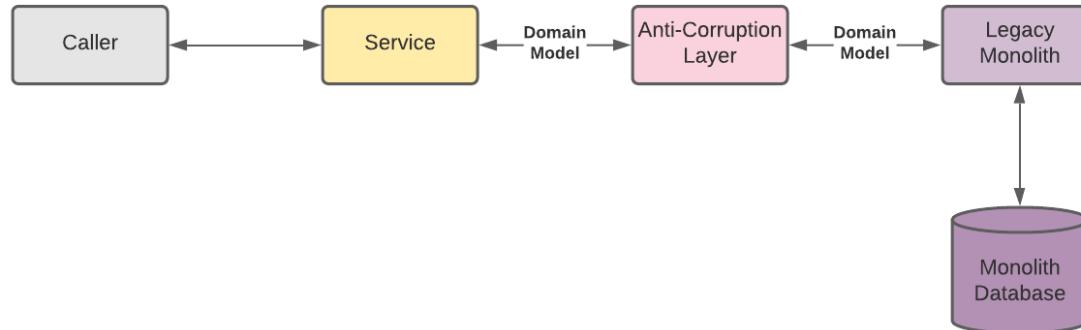
- Helps to optimize network traffic by preventing calls that are going to fail
- Can help prevent “noise” and network overload

## Potential “gotchas”:

- Process supported needs to be able to absorb limited downtime
- If not managed correctly, can result in poor user experience

# Microservice Design Patterns – Anti-Corruption Layer

- Prevents corruption of new microservice domain model(s) with legacy monolith domain model(s)
- Provides a proxy/translation layer to help keep models “clean”



# Microservice Design Patterns – Anti-Corruption Layer

## Benefits:

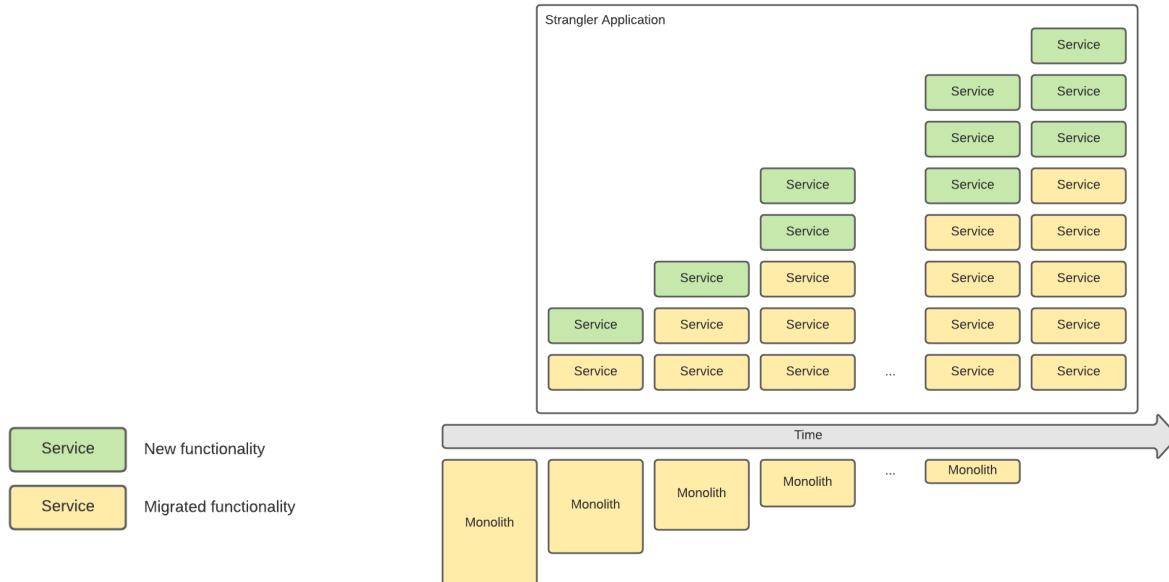
- Prevents crossing of boundaries and “leaking” of sub-optimal logic
- Helps to keep this insulation transparent between communicating parties

## Potential “gotchas”:

- Additional complexity
- Additional testing required to validate the additional complexity

# Microservice Design Patterns – Strangler Application

- Supports breaking up a monolith into microservices over time
- Can accommodate existing as well as new functionality



# Microservice Design Patterns – Strangler Application

Benefits:

- Allows a gradual vs. “big bang” breakup
- Especially well-suited for large monoliths that are taking active traffic

Potential “gotchas”:

- Requires business to be able to absorb gradual
- Adds complexity – coordinating new functionality, migrated functionality, and proper integration between them
- While transition in progress, will have to maintain two paths

## Group Discussion:

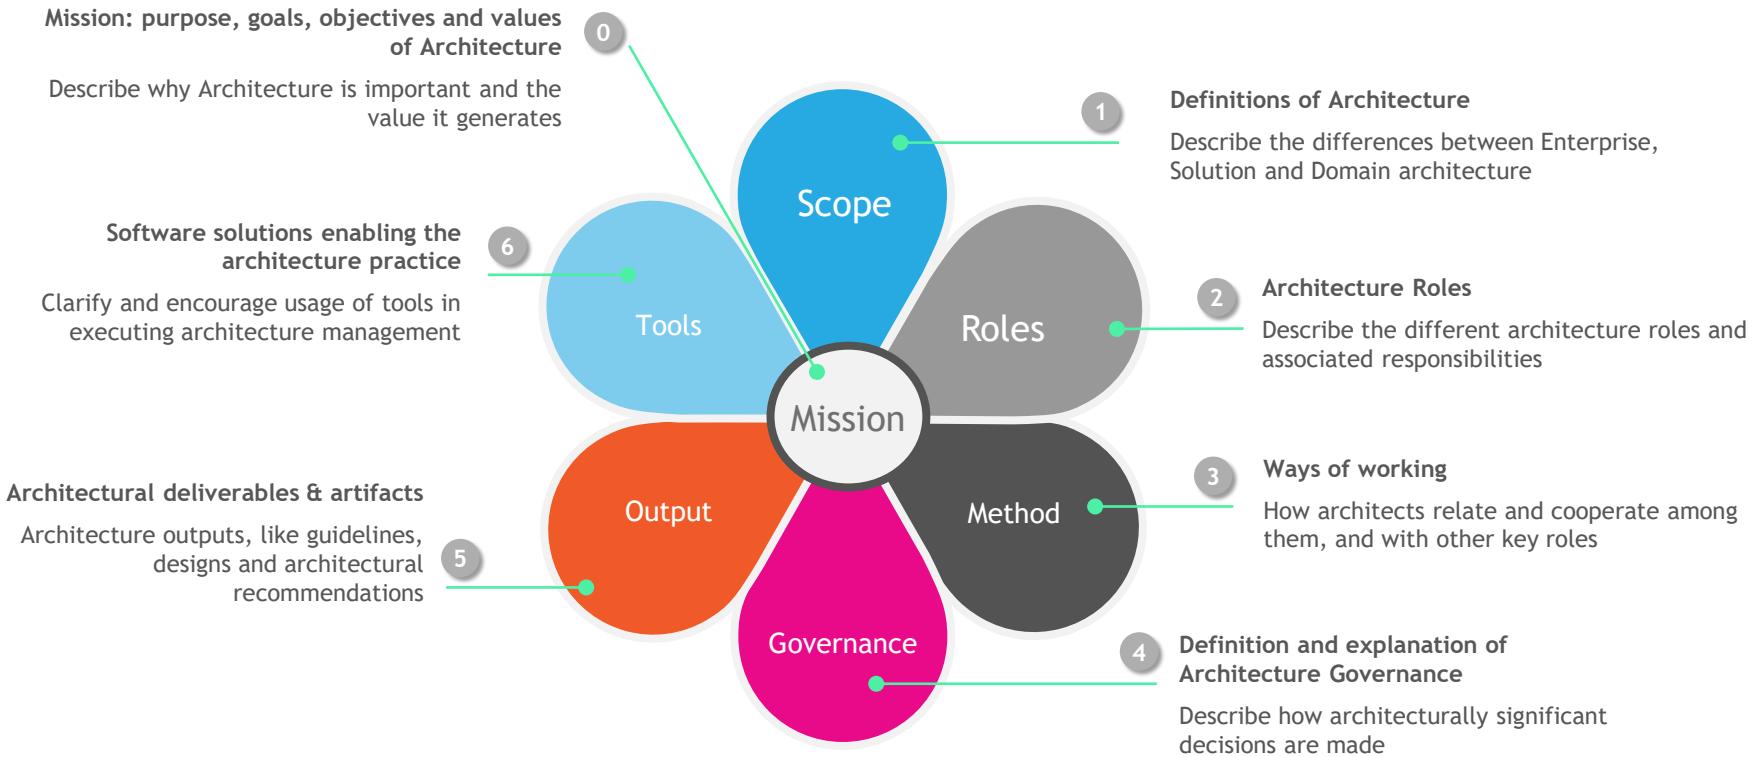
### Microservices Architecture

As a group, discuss the pros & cons of this architectural style.  
Specifically:

- What are some advantages to using this style? Is there a specific application profile that would cause you to prefer this style?
- What are some disadvantages to using this style? Is there a specific application profile that would cause you to avoid this style?

# Solution & Evolutionary Architecture

# Understanding 7 dimensions of architecture



# The Mission



Architecture seeks to build out capabilities that enable the business vision of the organization through technology



Architecture seeks to avoid different (sometimes substandard) approaches to solving the same business or technical problem because that has a high cost & high risk



Architecture seeks to help an organization minimize “accidental complexity”



Architecture seeks to provide guidance to the larger organization on standards & best practices for optimal project execution



Architecture, in concert with other teams, seeks to provide oversight of an organization’s technology landscape to ensure that it improves rather than degrades over time

# Solution Architecture vs. Application Architecture

TOGAF defines the terms “Solution Architecture”, and “Application Architecture” as follows:

## 3.69 Solution Architecture

A description of a discrete and focused business operation or activity and how IS/IT supports that operation.

**Note:**

A Solution Architecture typically applies to a single project or project release, assisting in the translation of requirements into a solution vision, high-level business and/or IT system specifications, and a portfolio of implementation tasks.

## 3.3 Application Architecture

A description of the structure and interaction of the applications as groups of capabilities that provide key business functions and manage the data assets.

Source: <https://pubs.opengroup.org/architecture/togaf9-doc/arch/chap03.html>

# Solution Architecture vs. Application Architecture



- Often described as a function of scope
- Am I focused on one or more specific solutions or services? (Solution Architecture)
- Or, am I focused on the discipline of architecting applications to enable business & technical capabilities? (Application Architecture)

In alignment with the other types of architectural roles, the goal is consistency and coordination

- Can we use a specific project to move the organization forward architecturally?
- Does the architecture for this solution (or this capability) align with our enterprise standards & best practices?

# Architecture Design in Practice



# Methodology

# Methodology



What is the “list” of concerns we should be focused on to ensure organizational success?



How do we effectively coordinate those multiple architectural concerns across the enterprise?



Is there a proven & repeatable process that we can follow to ensure we are documenting & managing to the right areas of concern?



Does this process support development of the architectural capability over time as our organization matures?

# Enter TOGAF



TOGAF  
SERIES  
GUIDES



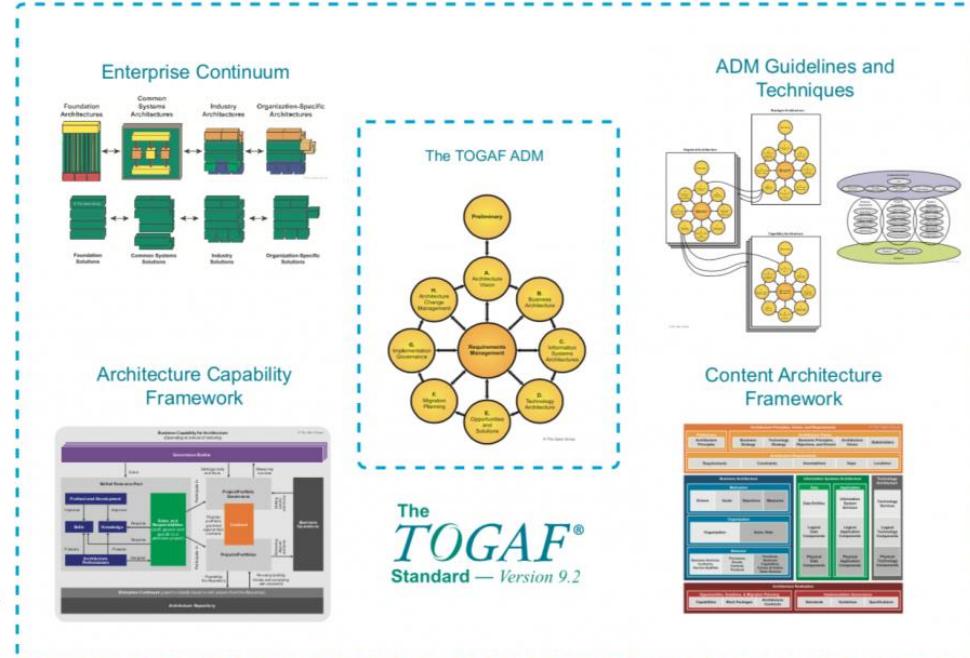
STANDARDS



BLOGS



CERTIFICATIONS



REFERENCE  
ARCHITECTURES



POCKET  
GUIDES



REFERENCE  
CARDS



DATA  
SHEETS



CASE  
STUDIES

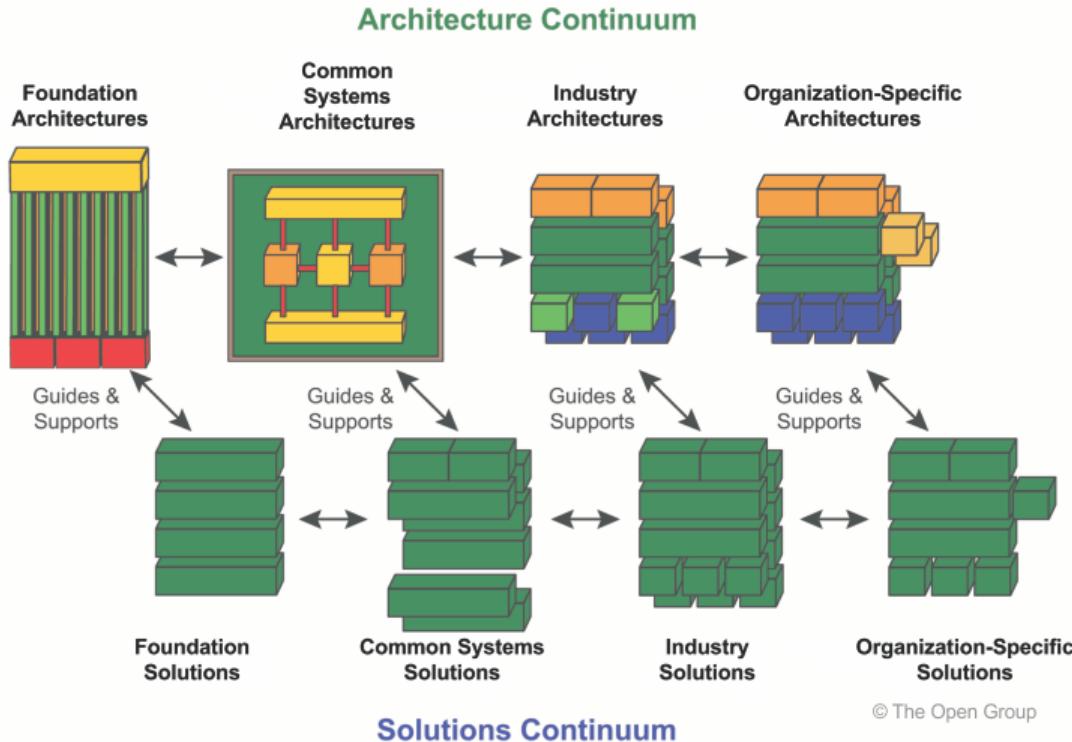


WEBINARS



WHITE  
PAPERS

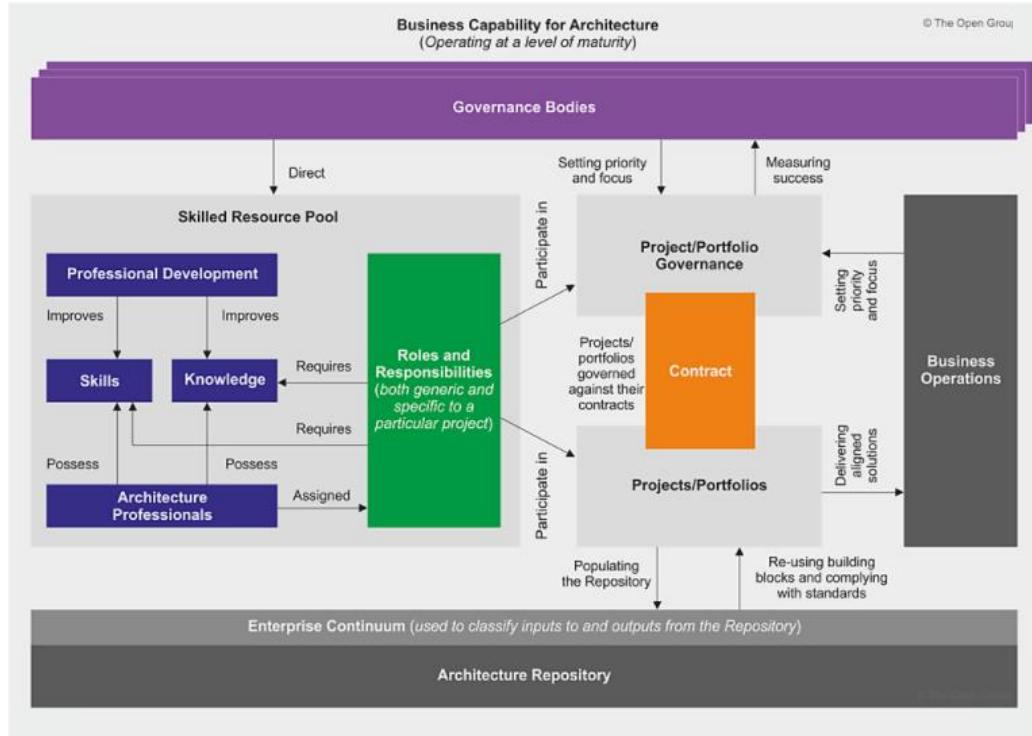
# Enterprise Continuum



- Enterprise Continuum is made up of Architecture and Solutions Continuums
- Provides context for the full set of architecture & solution assets that apply across the enterprise
- Architecture Continuum classifies architecturally significant assets (from generic to specific)
- Solutions Continuum classifies specific solutions that implement and achieve architectural intent

© The Open Group

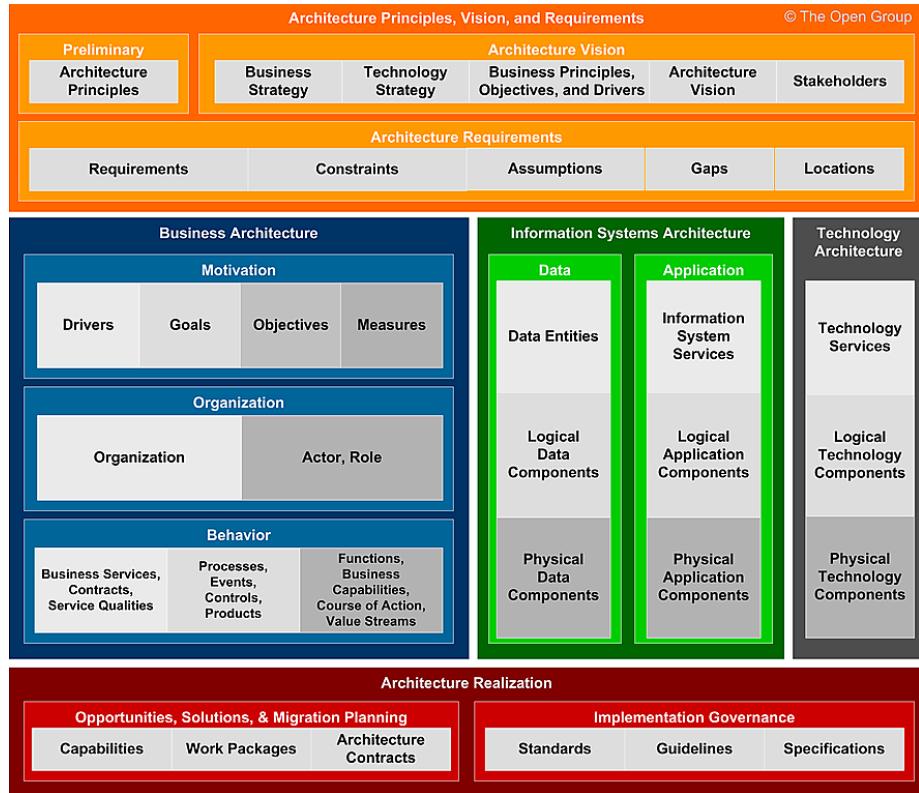
# Architecture Capability Framework



Defines the organizational structures, processes, roles, responsibilities, and skills/knowledge needed to realize a mature architectural capability (ensuring business alignment)

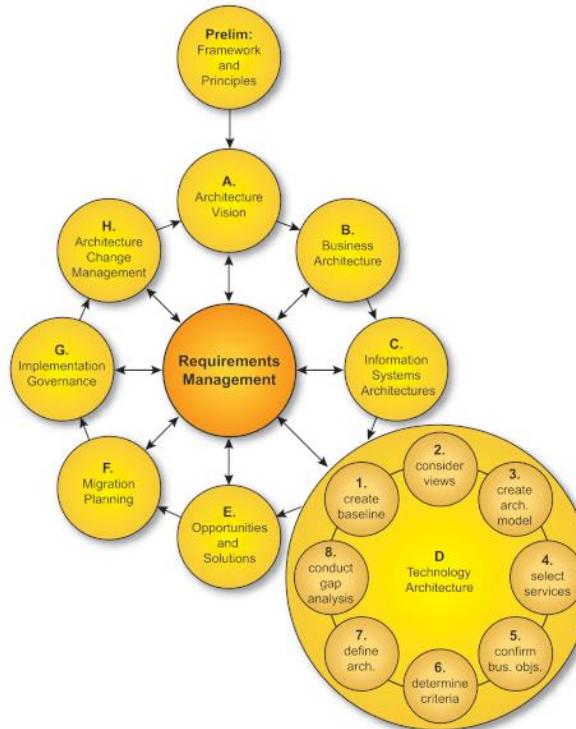
# Architecture Content Framework

Provides a structured taxonomy for the types of architectural artifacts and content that can prove useful to an enterprise when building out an architectural capability



# Architecture Development Method (ADM)

- Describes an iterative methodology for defining and building an enterprise architecture capability
- Reflects cyclical nature of continuous feedback and improvement
- Includes ongoing assessment of progress and activity at each step against defined requirements
- Each step further outlines a set of sub-steps that inform specific details of an activity
- Structured and meant to be adapted to your organization



# Use Case – Methodology

As a squad, you would like to “sell” to the business a modernization effort to migrate a **monolithic, mainframe-based application** used to manage a **critical sales & distribution process** within the company to a **microservices** architecture. Included in the target goals would be the transition of the workflow from the data center to **Cloud and Cloud native technologies** (where possible).

Currently, the application leverages a mainframe-based relational data store to store and serve all data within the large and tightly-coupled workflow. There are several legacy technologies in play and the platform upon which the application is deployed is **several versions behind latest**. Key reports that the business needs for critical decisioning exist but **cannot be delivered in real-time**. They are requested in batch and sent as a follow-up, sometimes 1 or 2 days later.

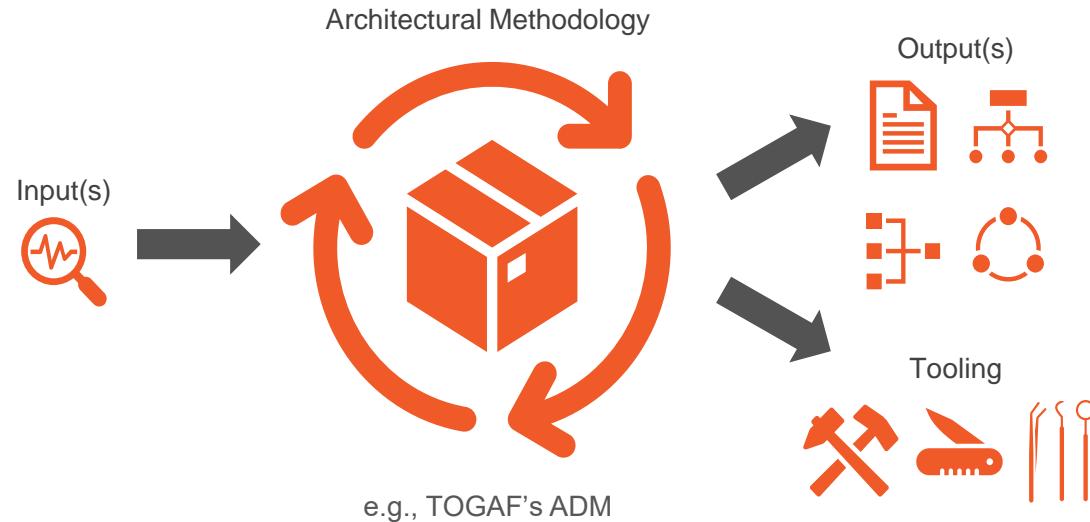
Also, the business would like to be able to leverage the different dimensions of data housed within this system for **Machine Learning** purposes but, at present, the data is difficult to extract and utilize. The current mainframe system supports **external customer access through a Web portal, internal access for account & order management (including remote sales associates utilizing a mobile device)**, and a handful of API endpoints that are used by **external partners** to enable **authorized third-party sales**.

In this scenario:

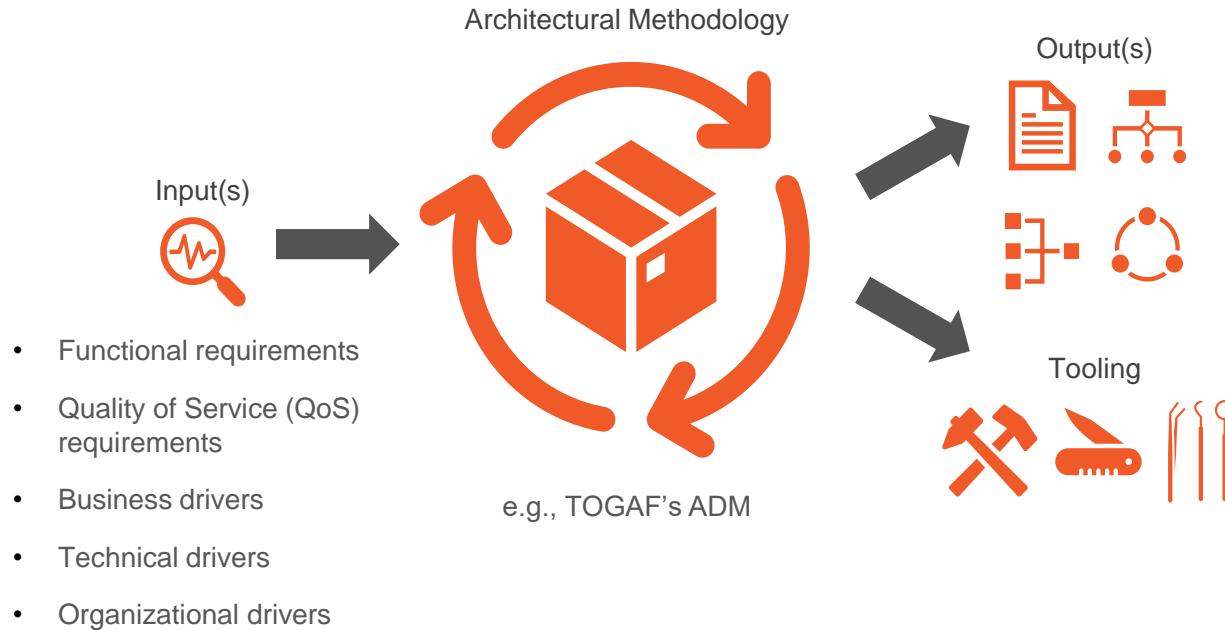
- What do you see as the business, technology, data, and application architecture concerns?

# Outputs & Tooling

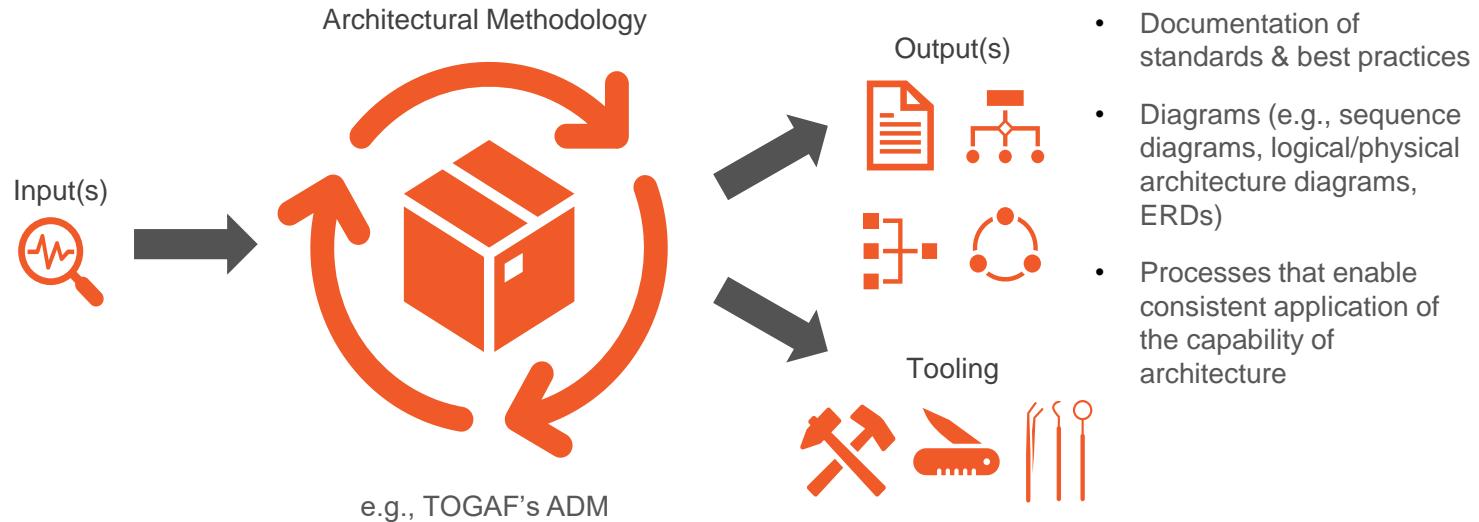
# Architectural Artifacts



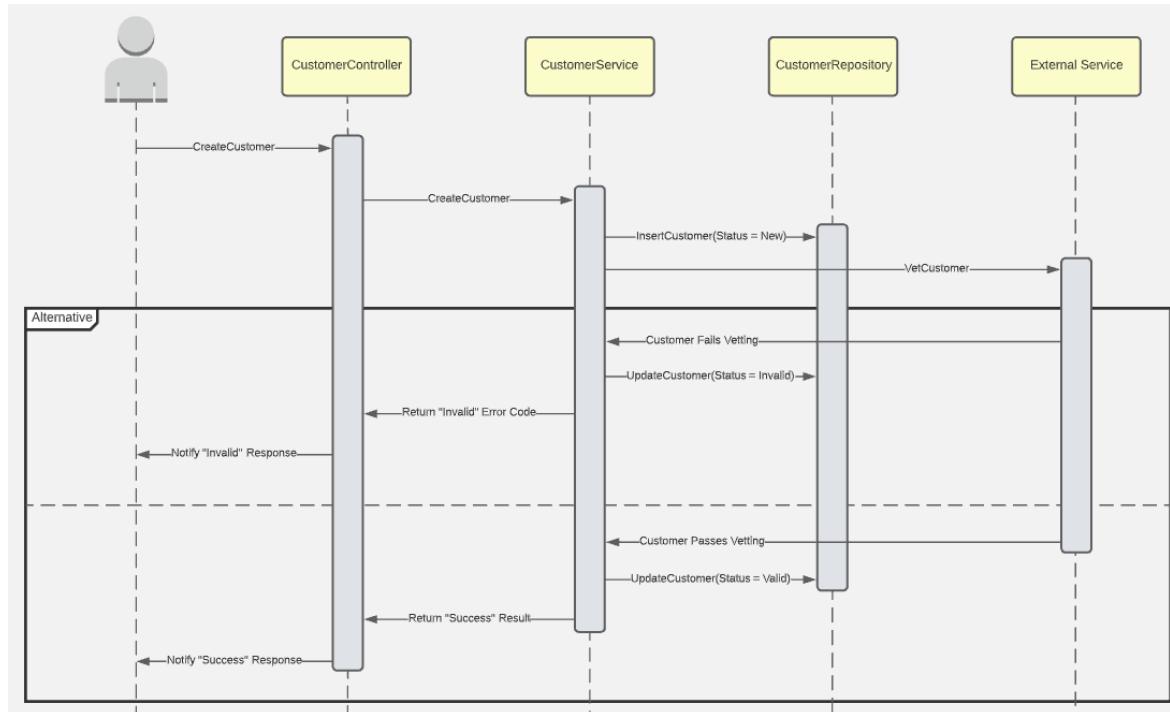
# Architectural Artifacts



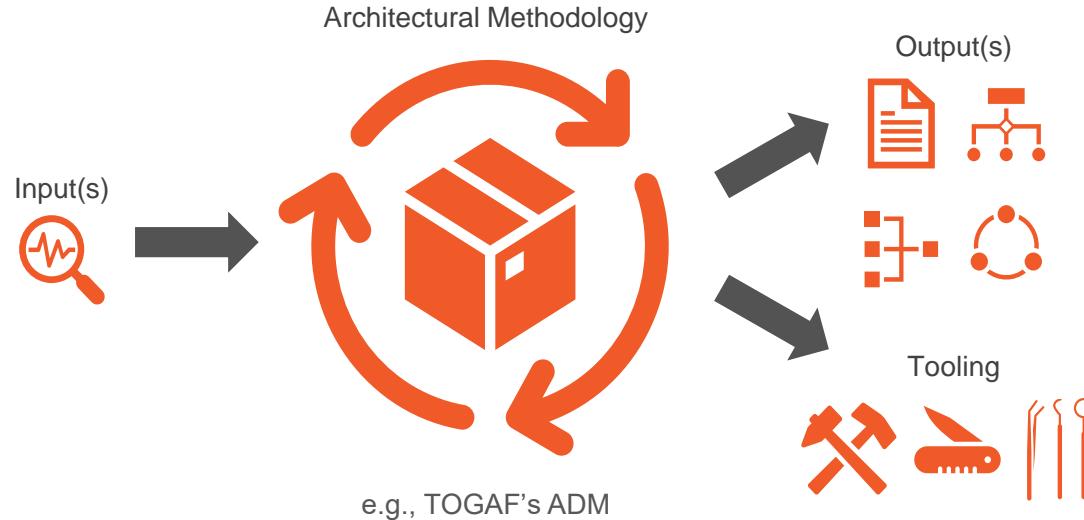
# Architectural Artifacts



# Sample Sequence Diagram

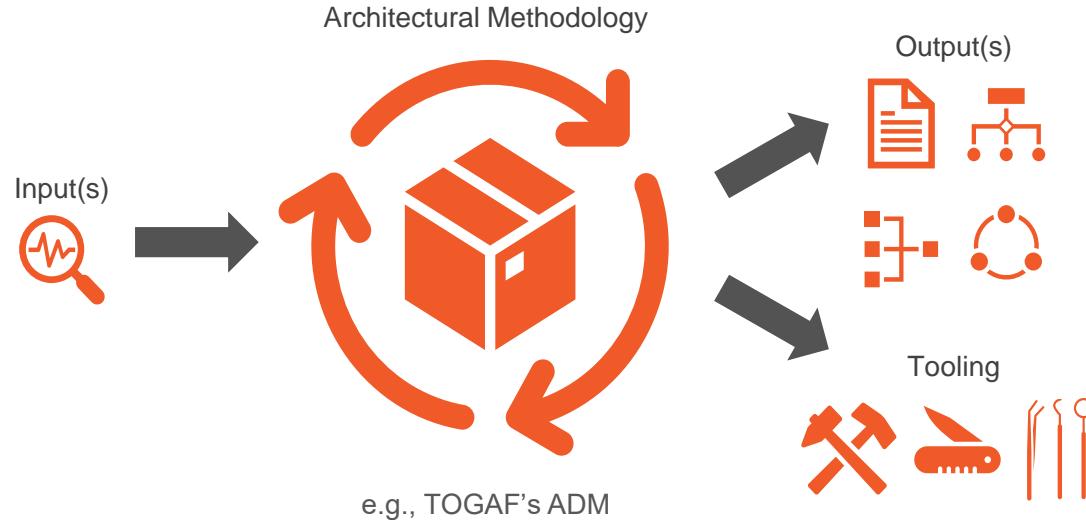


# Architectural Artifacts



- SDKs, frameworks, libraries, etc. that can be integrated into new projects & designs
- Helps promote consistency & speed
- Making the right thing the “easy” thing

# Architectural Artifacts



Couple of key points to keep in mind:

- Output & tooling may vary depending on architectural scope (enterprise, solution, domain), and all scopes should be considered
- Outputs & tooling must be easy to locate (search capabilities, known & maintained repositories, etc.); the architectural methodology should assist with cataloguing

# Use Case – Outputs & Tooling

As a squad, you would like to “sell” to the business a modernization effort to migrate a **monolithic, mainframe-based application** used to manage a **critical sales & distribution process** within the company to a **microservices** architecture. Included in the target goals would be the transition of the workflow from the data center to **Cloud and Cloud native technologies** (where possible).

Currently, the application leverages a mainframe-based relational data store to store and serve all data within the large and tightly-coupled workflow. There are several legacy technologies in play and the platform upon which the application is deployed is **several versions behind latest**. Key reports that the business needs for critical decisioning exist but **cannot be delivered in real-time**. They are requested in batch and sent as a follow-up, sometimes 1 or 2 days later.

Also, the business would like to be able to leverage the different dimensions of data housed within this system for **Machine Learning** purposes but, at present, the data is difficult to extract and utilize. The current mainframe system supports **external customer access through a Web portal**, **internal access for account & order management (including remote sales associates utilizing a mobile device)**, and a handful of API endpoints that are used by **external partners** to enable **authorized third-party sales**.

In this scenario:

- What types of architecture artifacts would you anticipate producing? How might this project use them? How might future projects use them?
- What types of reusable tooling might arise out of the execution of this project effort?

# Quality of Service Requirements

# Quality attributes as Architecture Requirements

What are some practical ways to architect for these key QoS requirements?

Five fundamental quality attributes



## Scalability

Ability to address increasing and fluctuating volumes with acceptable response time



## Availability/Reliability

Ability to ensure high availability and recover from failures



## Flexibility

Ability to add, modify and remove features and capabilities



## Operability

Ability to support smooth operations and easy maintenance



## Security

Ability to secure the enterprise assets from internal and external threats

# Quality attributes as Architecture Requirements

- Includes a couple of options – scaling up & scaling out (scaling out most common type with microservices architectures)
- Statelessness is a useful paradigm
- Scaling rules can be resource-based or schedule-based
- Goal should be mechanisms that automate scaling out & scaling in
- Scaling in also important for cost management

Five fundamental quality attributes



## Scalability

Ability to address increasing and fluctuating volumes with acceptable response time



## Availability/Reliability

Ability to ensure high availability and recover from failures



## Flexibility

Ability to add, modify and remove features and capabilities



## Operability

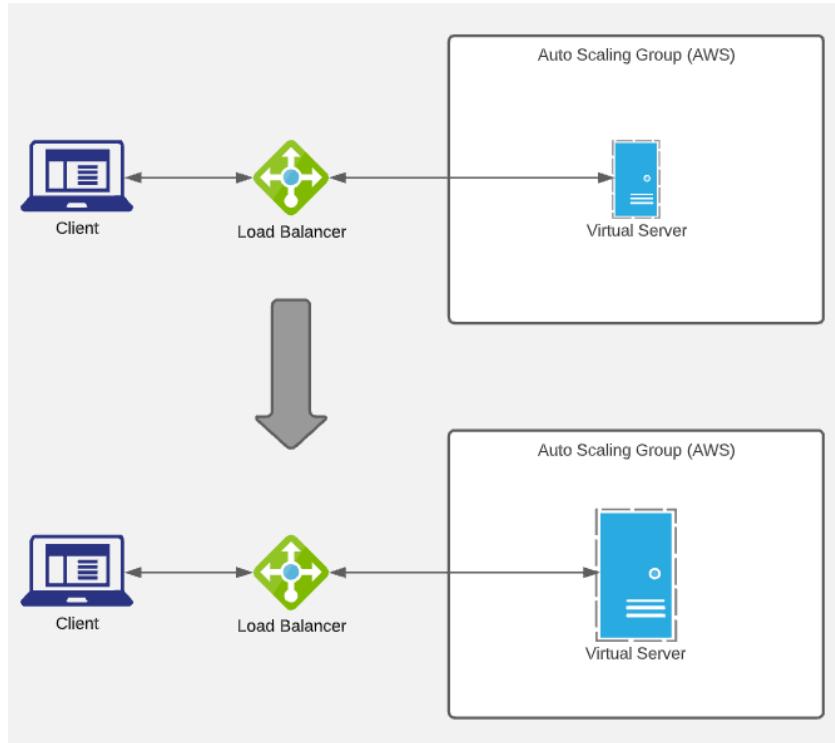
Ability to support smooth operations and easy maintenance



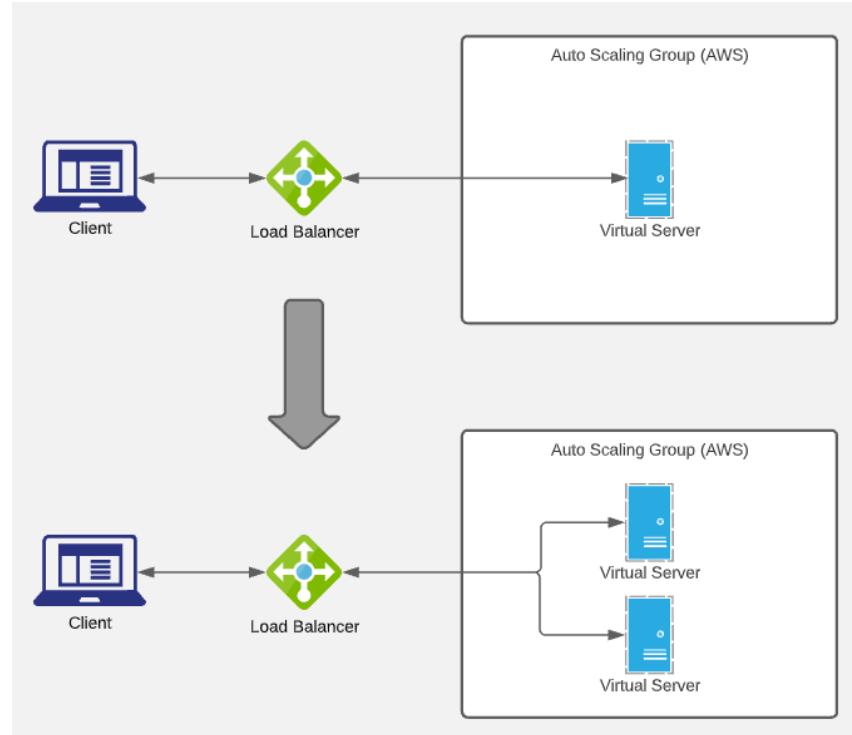
## Security

Ability to secure the enterprise assets from internal and external threats

# Vertical Scaling (Scaling “Up”)



# Horizontal Scaling (Scaling “Out”)



## Moment of Reflection



What are some pros & cons of each type of scaling  
(scaling “up” vs. scaling “out”)?

Type your answer in chat or come off mute & share

# Quality attributes as Architecture Requirements

- Need a clear understanding of minimum SLAs
- What are the requirements for geographic redundancy? Across the country or across the globe?
- Management of latency and distance data has to travel
- If transactional, is active-passive sufficient or do you require active-active?
- Recovery Time Object (RTO) & Recovery Point Objective (RPO) are important data points
- Cost optimization

Five fundamental quality attributes



## Scalability

Ability to address increasing and fluctuating volumes with acceptable response time



## Availability/Reliability

Ability to ensure high availability and recover from failures



## Flexibility

Ability to add, modify and remove features and capabilities



## Operability

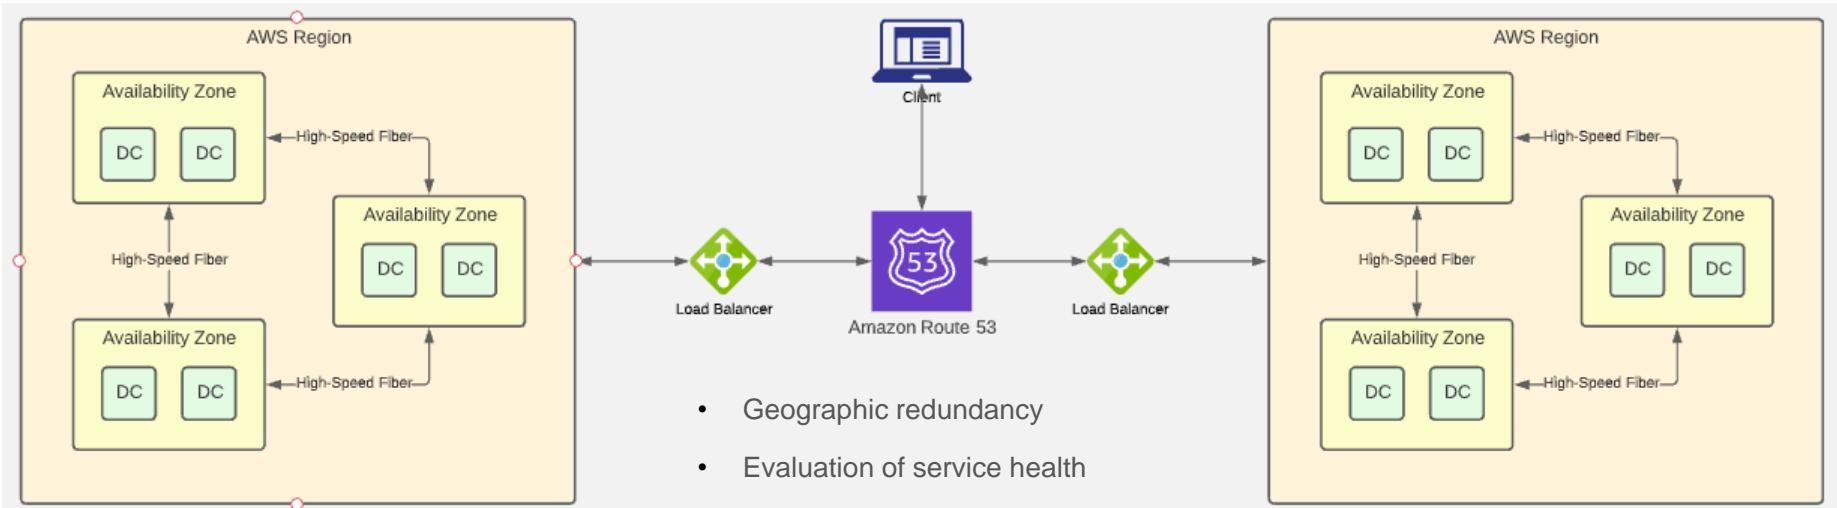
Ability to support smooth operations and easy maintenance



## Security

Ability to secure the enterprise assets from internal and external threats

# Availability/Reliability



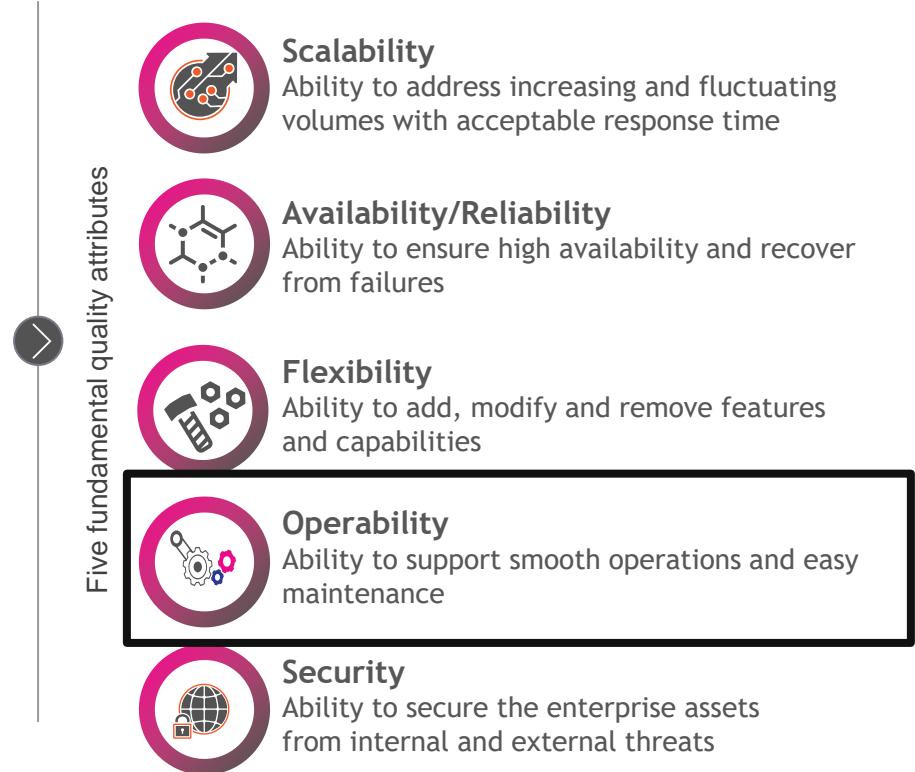
# Quality attributes as Architecture Requirements

- Ensuring that the architecture promotes loose coupling
- Leveraging design patterns that facilitate adding & removing features with minimized impact
- Maintaining a proper versioning strategy, especially if requirement to run multiple versions side-by-side
- Design to abstractions rather than concretions
- Utilize techniques like Dependency Injection
- Make maintainability a key goal of the design

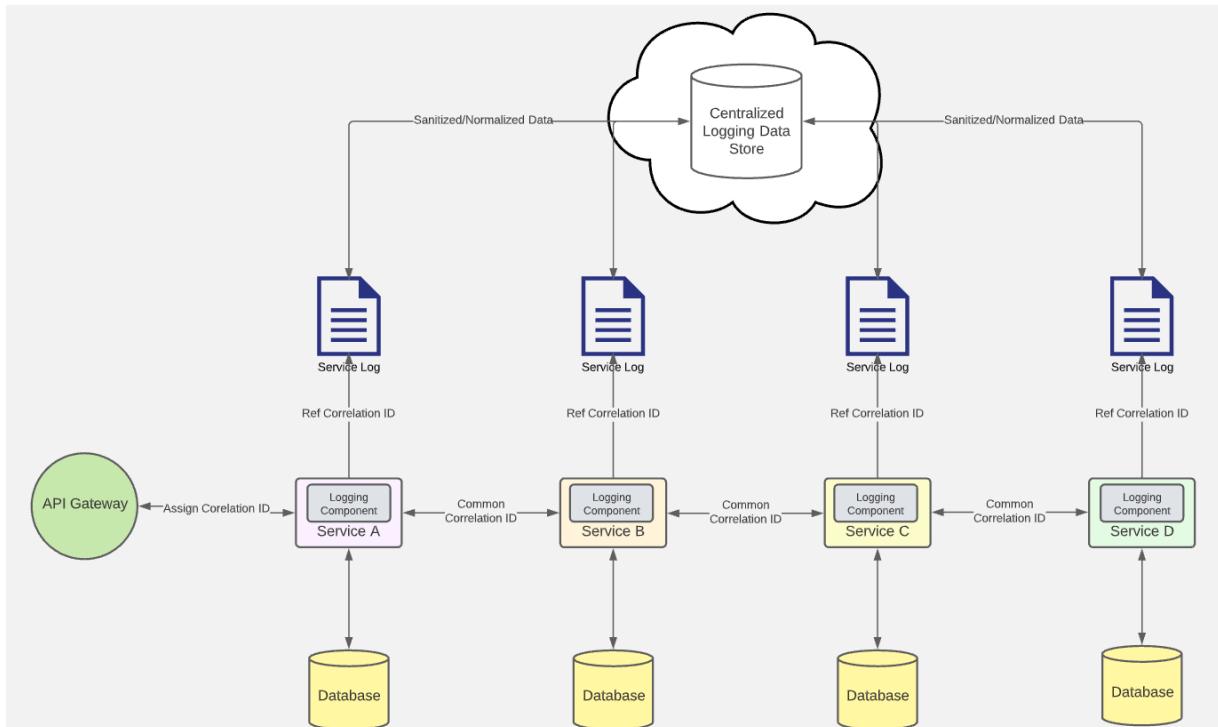


# Quality attributes as Architecture Requirements

- Use known observability patterns for logging & tracing transactions through the system
- Likely able to find multiple frameworks to enable – don't reinvent the wheel
- Ensure sanitization of relevant log data to prevent leakage of sensitive information
- Especially in a microservices architecture, end-to-end traceability becomes critical
- Correlation IDs can be transmitted across a series of microservice requests (e.g., in headers) to tie together
- Ask yourself "What would I need to know to support this system in production if I'm getting paged at 3 a.m.?"



# Operability



# Quality attributes as Architecture Requirements

- Practice key tenets of security architecture:
  - Defense in depth
  - Least privilege
  - Secure by default
- Utilize principles of Zero Trust Architecture to help you avoid “castle & moat”, minimizing zones of implicit trust within the system & its components
- Ideally, the architecture of the system will be built around an established DevSecOps framework which seeks to “shift security left”

Five fundamental quality attributes



## Scalability

Ability to address increasing and fluctuating volumes with acceptable response time



## Availability/Reliability

Ability to ensure high availability and recover from failures



## Flexibility

Ability to add, modify and remove features and capabilities



## Operability

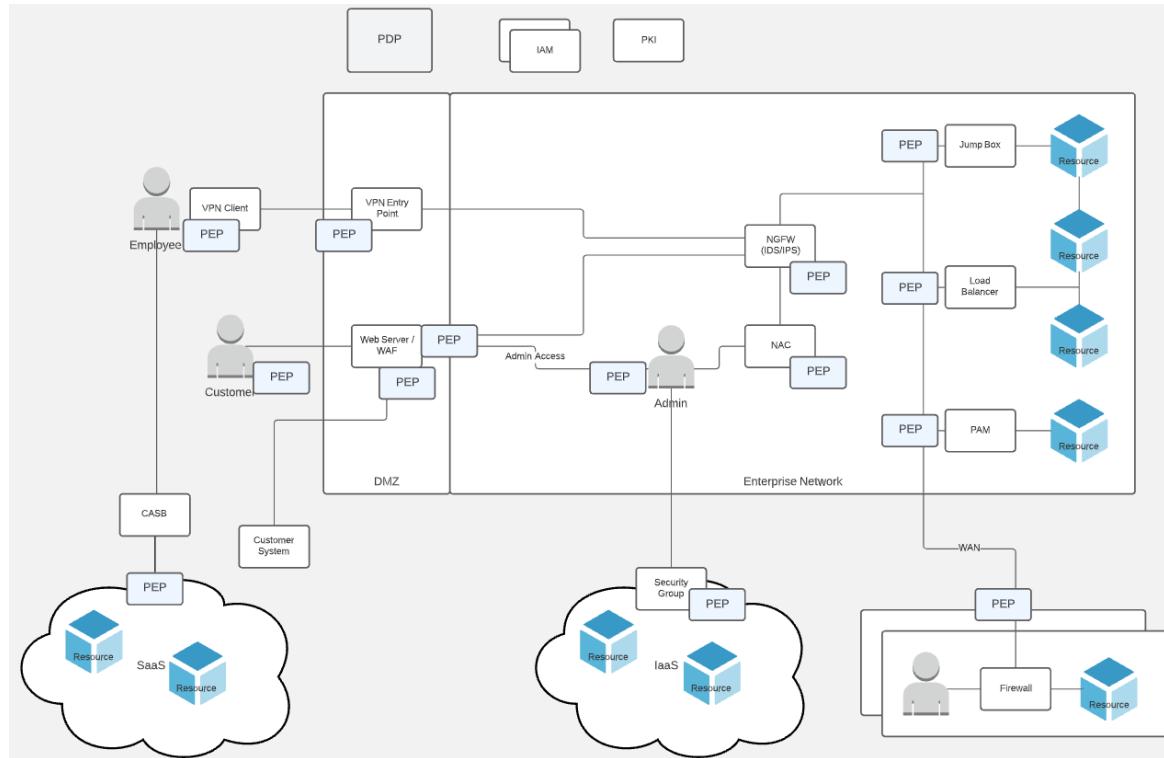
Ability to support smooth operations and easy maintenance



## Security

Ability to secure the enterprise assets from internal and external threats

# Security – Zero Trust Architecture



# Use Case – QoS Requirements

As a squad, you would like to “sell” to the business a modernization effort to migrate a **monolithic, mainframe-based application** used to manage a **critical sales & distribution process** within the company to a **microservices** architecture. Included in the target goals would be the transition of the workflow from the data center to **Cloud and Cloud native technologies** (where possible).

Currently, the application leverages a mainframe-based relational data store to store and serve all data within the large and tightly-coupled workflow. There are several legacy technologies in play and the platform upon which the application is deployed is **several versions behind latest**. Key reports that the business needs for critical decisioning exist but **cannot be delivered in real-time**. They are requested in batch and sent as a follow-up, sometimes 1 or 2 days later.

Also, the business would like to be able to leverage the different dimensions of data housed within this system for **Machine Learning** purposes but, at present, the data is difficult to extract and utilize. The current mainframe system supports **external customer access through a Web portal**, **internal access for account & order management (including remote sales associates utilizing a mobile device)**, and a handful of API endpoints that are used by **external partners** to enable **authorized third-party sales**.

In this scenario:

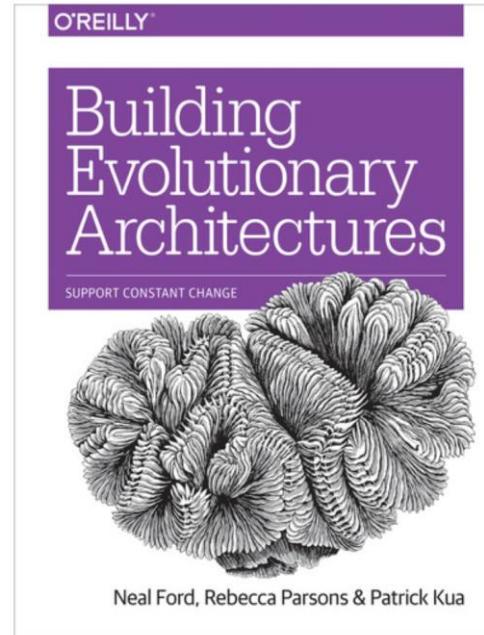
- How would you prioritize the 5 key QoS requirements for this use case?
- What are some practical ways you could architect for each?
- How might you measure adherence (or discover deficiency)?



# Evolutionary Architecture

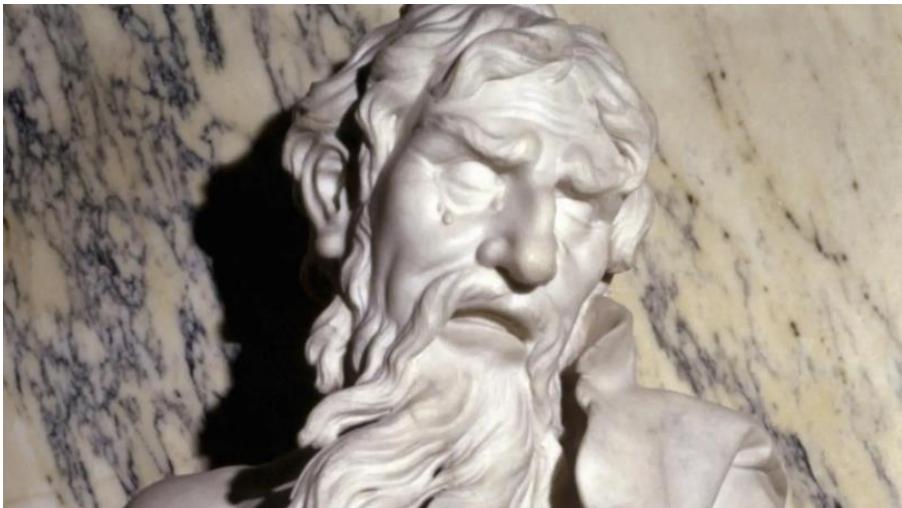


# Evolutionary Architecture



Great book on this topic!

# Changes

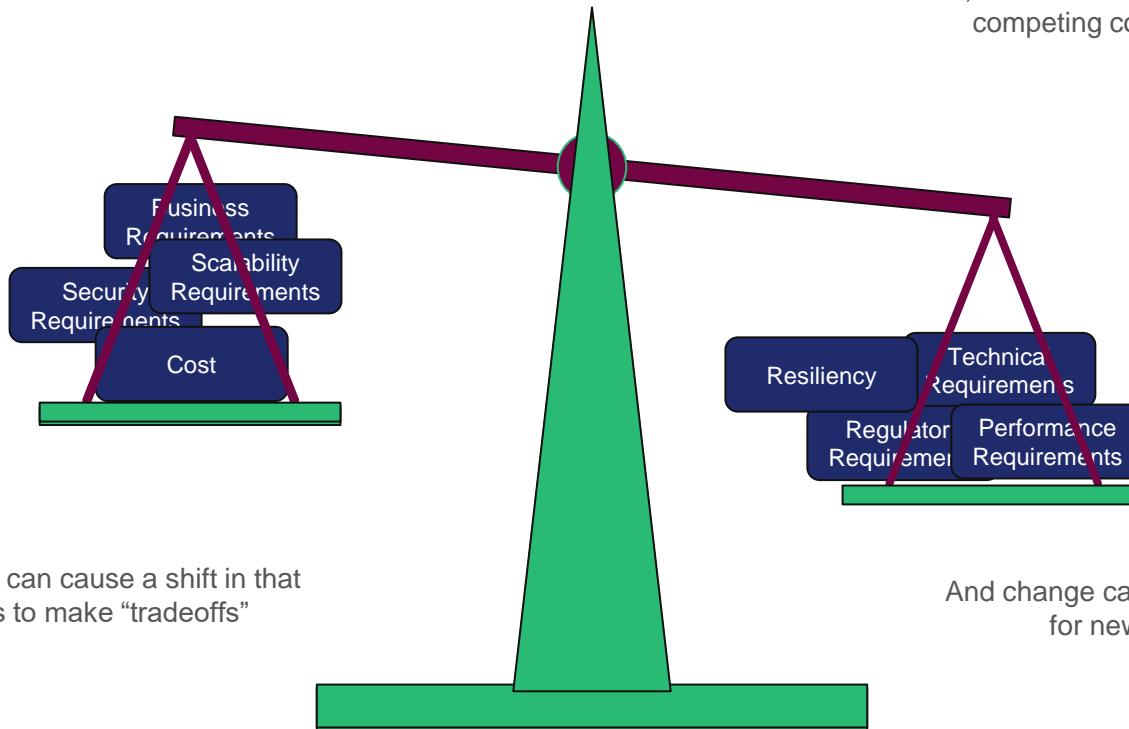


- Heraclitus (Greek philosopher) was quoted as saying “the only constant is change”
- When building software, we execute a project at a point in time
- However, the goal is for the software we build to provide value that outlasts the immediate
- For that to be a reality, we have to expect (and ideally architect for) changes
- Change comes in many forms

Change due to business innovation, mergers/acquisitions, a new regulatory requirement, evolving security threats, advancements in technology, etc.

# Architectural Balance

As architects, we are forced to balance multiple competing concerns

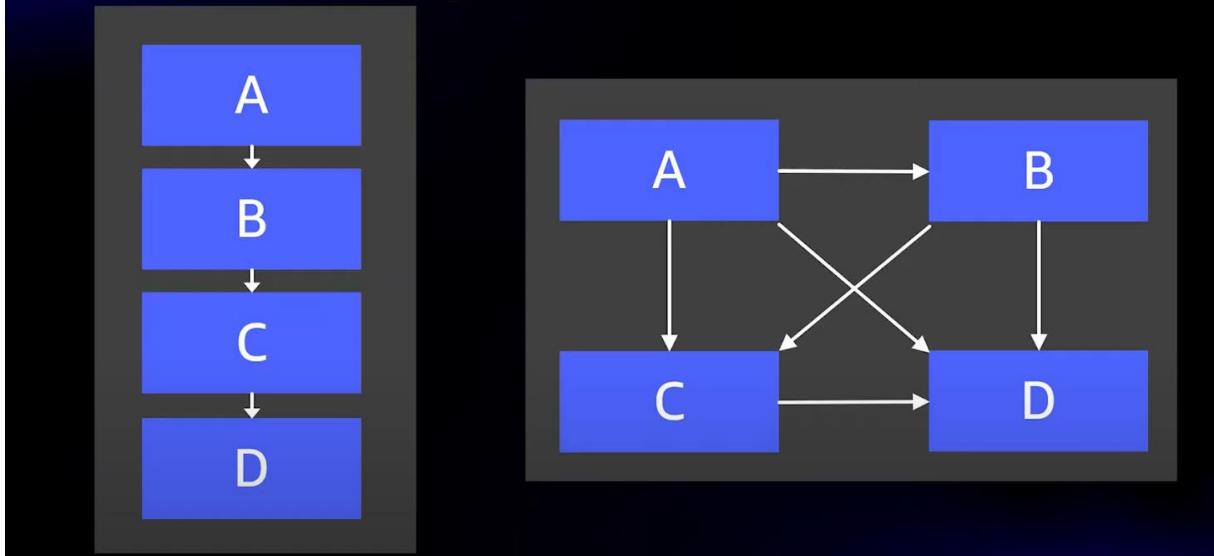


Each new requirement can cause a shift in that balance, forcing us to make “tradeoffs”

And change can be a common source for new requirements

# Balancing Competing Architectural Concerns

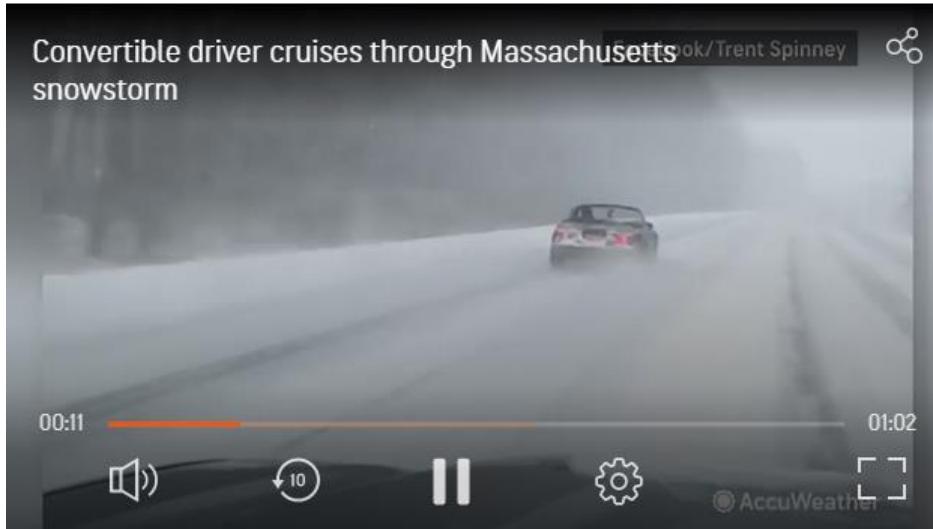
## Two system designs



Given these two options for designing a distributed system, what are the architectural tradeoffs for each? How would you balance the architectural concerns associated to each?

# Architectural Awareness

Balancing competing architectural concerns requires that we maintain awareness of our surroundings



Source: [Accuweather.com](https://www.accuweather.com)

We must be able to see both the  and the 

We cannot bury our heads in the sand (or snow); otherwise, we run the risk of missing critical details that can mean the difference between success and suboptimal

## Two Types of Change



### Incremental

How do we move our architectures forward in small but meaningful steps, balancing the shifts in equilibrium that can come with each change?



### Guided

What are the most important architectural characteristics (or dimensions) of our designs, how do we define progress, and how do we assess current state against target state?

## Incremental Change



- Can involve both how we build software and how we deploy it
- With incremental steps, smaller scope of change
- What is “far enough” but “not too far”?
- From a deployment perspective, can potentially run multiple versions side-by-side (e.g., for a microservice), allowing callers to gradually migrate from old to new
- When old is no longer used, can be deadfiled

The challenge here can be ensuring our technology landscape remains streamlined and as “uncluttered” as possible

# Guided Change



Prioritize the most important architectural characteristics of the design



Identify a target state for the system



Leverage objective mechanisms to assess progress toward target state and adherence to those prioritized characteristics



If we feel our system veering off-course from the defined objectives, course correct



This assessment coupled with the incremental nature of change can help prevent more significant issues



# Architectural Dimensions

- Includes the multiple “-ilities” typically associated with QoS requirements
- But can include other architectural concerns as well
- For example:



Technical

- Frameworks, libraries,  
languages



Data

- Schemas, table design,  
indexing, performance tuning



Security

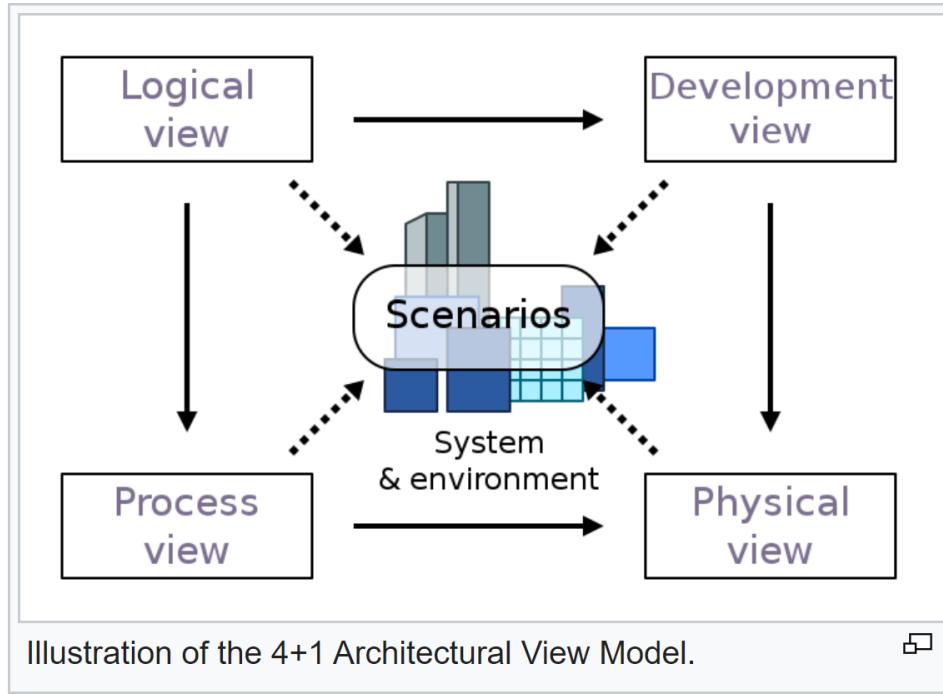
- Policies, guidelines, tooling,  
data protection



Operational/System

- How the architecture maps  
to existing infrastructure

# 4+1 Architectural View Model

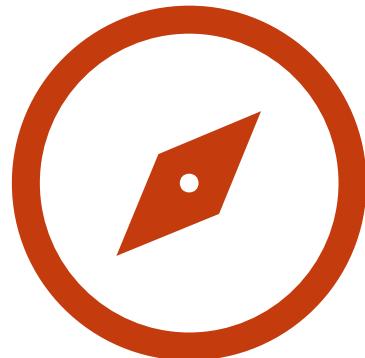


- Logical view – the functionality provided by the system to end-users
- Process view – the dynamic, runtime behaviour of the system; the system workflow
- Development view – the programmer's perspective; AKA the implementation view
- Physical view – the deployment view; what does the system look like running on production servers and production networks
- Scenarios – use cases that provide validation and inform the architecture

# Fitness Functions



- As part of guided change, fitness functions are those objective measures we can use to assess and evaluate progress
- Multiple types can be leveraged for different prioritized architectural dimensions
- Methodology of execution can vary as well depending on the dimension in question
- Ideally, provide a “quality gate” that can be used to identify issues early and prevent sub-optimal results



# Fitness Functions



## Atomic

- Exercise and verify a specific unit of architectural quality
- For example, a unit test that automatically validates a key architectural characteristic

## Holistic

- Exercise a combination of architectural concerns
- For example, performance & scalability
- Ensure combinations of architectural concerns “play well” together

# Fitness Functions



## Triggered

- Run based on the occurrence of a specific event
- For example, a build pipeline that ensures passing tests before deployment

## Continual

- Run continually to assess architectural quality over time
- For example, average transaction speed traced across a series of requests

# Fitness Functions



## Static

- Binary pass/fail

## Dynamic

- Nuanced evaluation
- For example, accounting for some acceptable level of performance hit as the system scales to handle a large volume of requests

# Fitness Functions



## Intentional

- What's known

## Emergent

- What's learned as design and development of the system proceeds

# Fitness Functions



## Examples

- Integrating combination of all unit tests passing with minimum coverage percentages into CI/CD pipeline, failing build if target measure unmet
- 0 critical or high security vulnerabilities found from static scans, dynamic scans, or open source component scans
- Successful adherence to RTO & RPO during disaster recovery exercises
- Adherence of average response times as measured against target SLAs under normal load (50<sup>th</sup>, 90<sup>th</sup>, 95<sup>th</sup> percentiles)
- Adherence of average response times as measured against target SLAs under stress (50<sup>th</sup>, 90<sup>th</sup>, 95<sup>th</sup> percentiles)

# Use Case – Fitness Functions

As a squad, you would like to “sell” to the business a modernization effort to migrate a **monolithic, mainframe-based application** used to manage a **critical sales & distribution process** within the company to a **microservices** architecture. Included in the target goals would be the transition of the workflow from the data center to **Cloud and Cloud native technologies** (where possible).

Currently, the application leverages a mainframe-based relational data store to store and serve all data within the large and tightly-coupled workflow. There are several legacy technologies in play and the platform upon which the application is deployed is **several versions behind latest**. Key reports that the business needs for critical decisioning exist but **cannot be delivered in real-time**. They are requested in batch and sent as a follow-up, sometimes 1 or 2 days later.

Also, the business would like to be able to leverage the different dimensions of data housed within this system for **Machine Learning** purposes but, at present, the data is difficult to extract and utilize. The current mainframe system supports **external customer access through a Web portal, internal access for account & order management (including remote sales associates utilizing a mobile device)**, and a handful of API endpoints that are used by **external partners** to enable **authorized third-party sales**.

In this scenario:

- Choose 1 of the Quality of Service requirements discussed previously. Can you think of a specific fitness function that could be used to help guide & validate the evolution of the architecture over time relative to that QoS requirement?
- Where (if anywhere) might you be able to apply incremental change in the execution of this overall effort?

-  **Scalability**  
Ability to address increasing and fluctuating volumes with acceptable response time
-  **Availability/Reliability**  
Ability to ensure high availability and recover from failures
-  **Flexibility**  
Ability to add, modify and remove features and capabilities
-  **Operability**  
Ability to support smooth operations and easy maintenance
-  **Security**  
Ability to secure the enterprise assets from internal and external threats

# Domain-Driven Design (DDD)

# Domain- Driven Design (DDD)

- Process of designing systems aligned to business domain
- Drives business language into the technology – engineer in terms of business process & impact
- Can enable IT to be an enabler and even an accelerator (vs. another “cost center”)
- Helps with management of types of coupling
- Direct correlation can be drawn between domains and microservices



# Defining Boundaries



- Enforcing strong boundaries and controlled interfaces can help mitigate coupling
- Domains become realms of cohesion
- Controlled integration – i.e., what do I allow others to “see”?
- DDD is about helping us determine where the boundaries should exist

# Some Key Terms

## Bounded Context

Provides a place for a set of terms to have a specific business meaning

Contexts allow for variance in technology from domain-to-domain

An entity may exist with the same name in two different bounded contexts, but its context will dictate its meaning



## Some Key Terms



Ubiquitous Language

- Shared language between tech and business

- Using language that models the business

- Promotes collaboration between domain experts and tech

# Some Key Terms

## Domain

A “slice” of your software system

Ideally, maps to a bounded context (and potentially one or more microservices)

Can encapsulate subdomains

Three major types: Generic, Supporting, and Core



## Some Key Terms

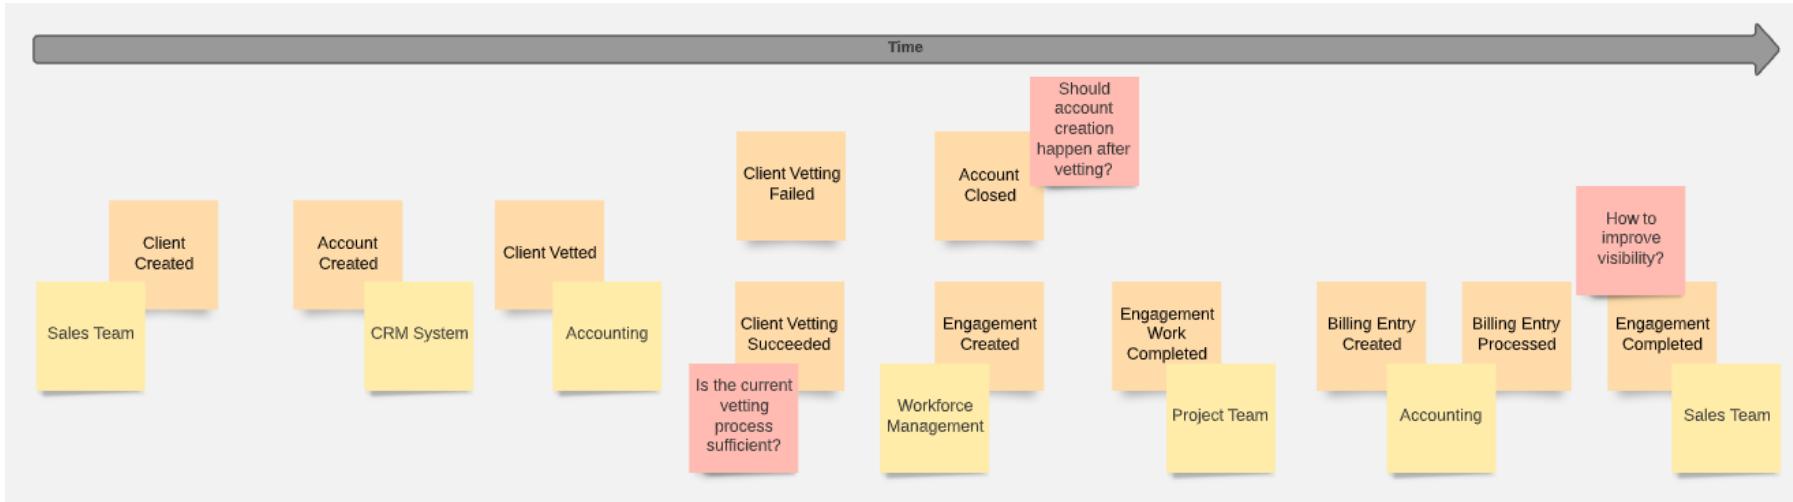
Generic Domain – “Buy vs. Build”, SaaS, use 3<sup>rd</sup> party or framework

Supporting Domain – Develop, maybe outsource

Core – Unique and distinguishing; where lion's share of efforts are focused



# EventStorming - Example



# DDD – Cheat Sheet



Strategic

Analyze Business Domain

Identify Subdomains  
Core, Supporting, Generic

Domain Modeling

Ubiquitous Language

Defining & Integrating  
Bounded Contexts

Tactical

Layered Architectures

Defining Services  
Domain & Application

Defining Aggregates

Defining Entities & Value Objects

Defining Domain Events

# Use Case – Architectural Styles & Patterns

As a squad, you would like to “sell” to the business a modernization effort to migrate a **monolithic, mainframe-based application** used to manage a **critical sales & distribution process** within the company to a **microservices** architecture. Included in the target goals would be the transition of the workflow from the data center to **Cloud and Cloud native technologies** (where possible).

Currently, the application leverages a mainframe-based relational data store to store and serve all data within the large and tightly-coupled workflow. There are several legacy technologies in play and the platform upon which the application is deployed is **several versions behind latest**. Key reports that the business needs for critical decisioning exist but **cannot be delivered in real-time**. They are requested in batch and sent as a follow-up, sometimes 1 or 2 days later.

Also, the business would like to be able to leverage the different dimensions of data housed within this system for **Machine Learning** purposes but, at present, the data is difficult to extract and utilize. The current mainframe system supports **external customer access through a Web portal**, **internal access for account & order management (including remote sales associates utilizing a mobile device)**, and a handful of API endpoints that are used by **external partners** to enable **authorized third-party sales**.

In this scenario:

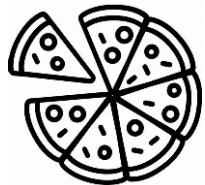
- First “blush”, where might you draw domain boundaries around the different areas of the system?
- Pick at least 2 of the previously discussed architectural styles (Layered, Microkernel, Event-Driven, & Microservices) – as a team, discuss high-level flow for an architecture to meet this use case that incorporates those two (or more) styles
- On return from group discussion, come prepared (using diagrams, notes, documentation, testing strategies, fitness functions, etc.) to share specific details of your selected architecture

# Architecting for the Cloud

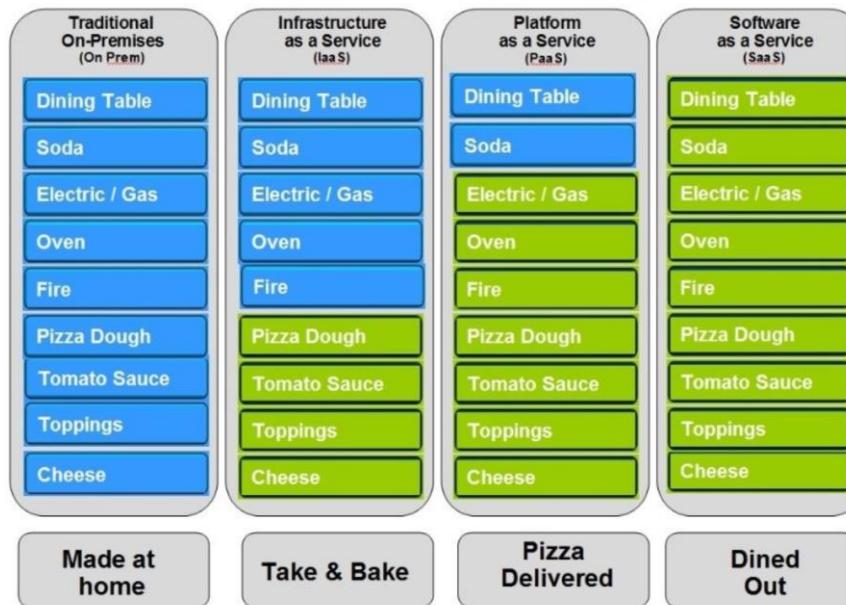


# Pizza-as-a-Service

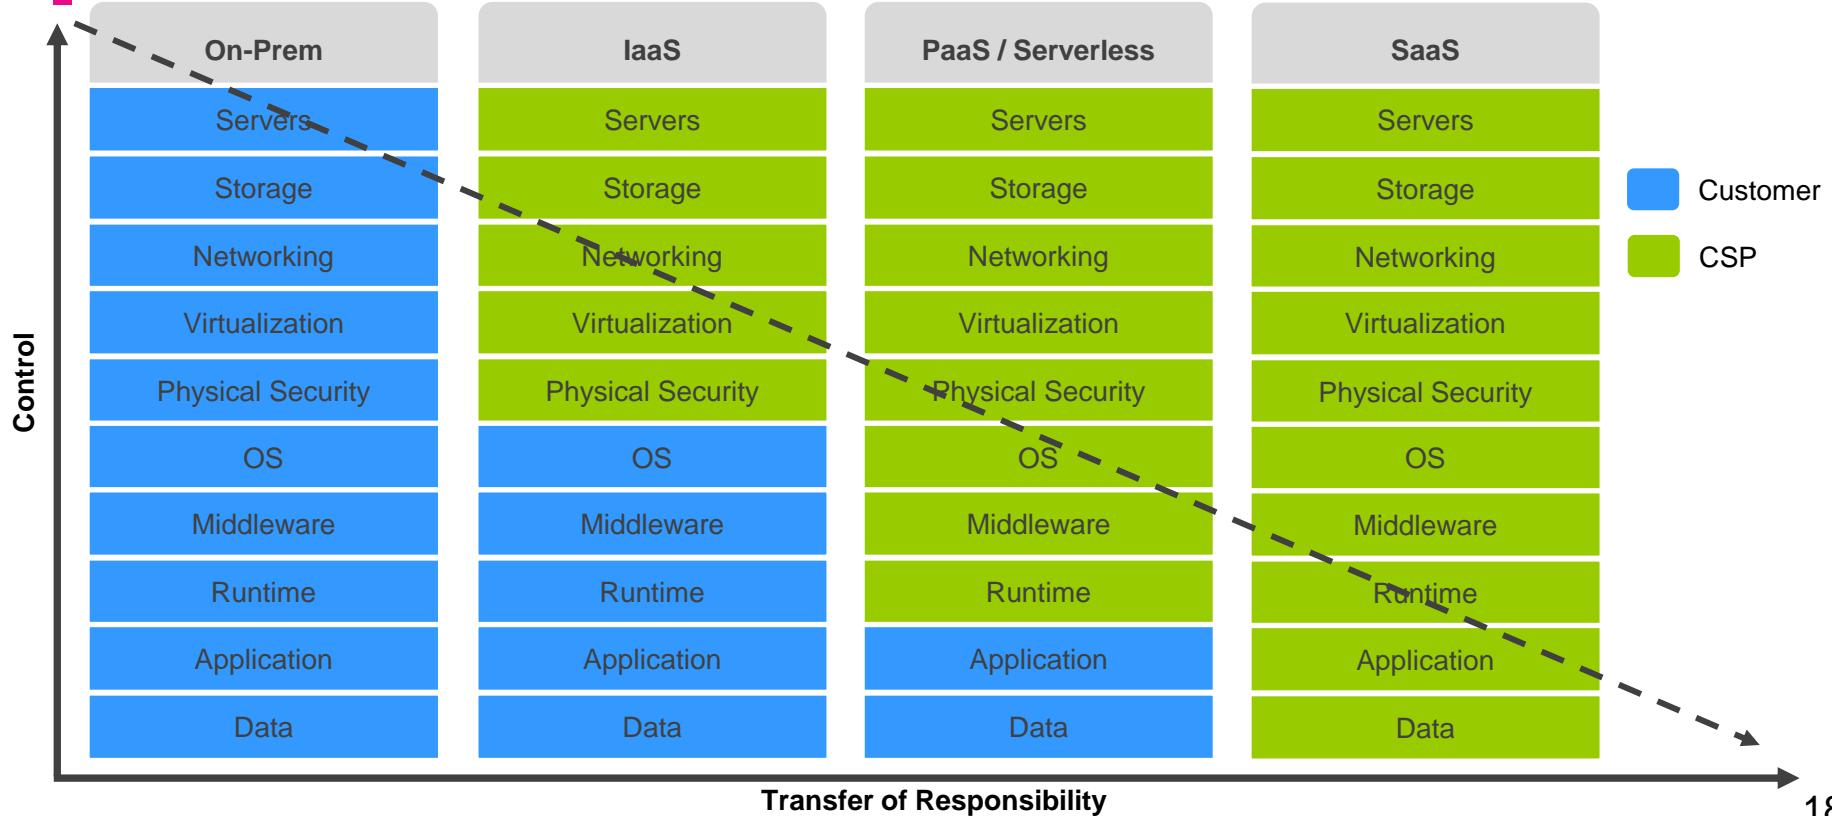
From a LinkedIn post by Albert Barron from IBM (<https://www.linkedin.com/pulse/20140730172610-9679881-pizza-as-a-service/>)



## Pizza as a Service

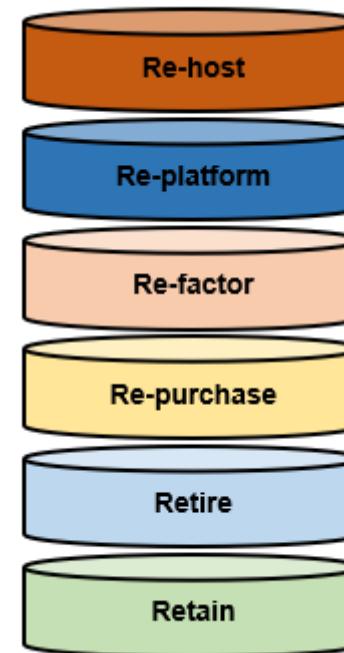


# Side-by-Side Comparison



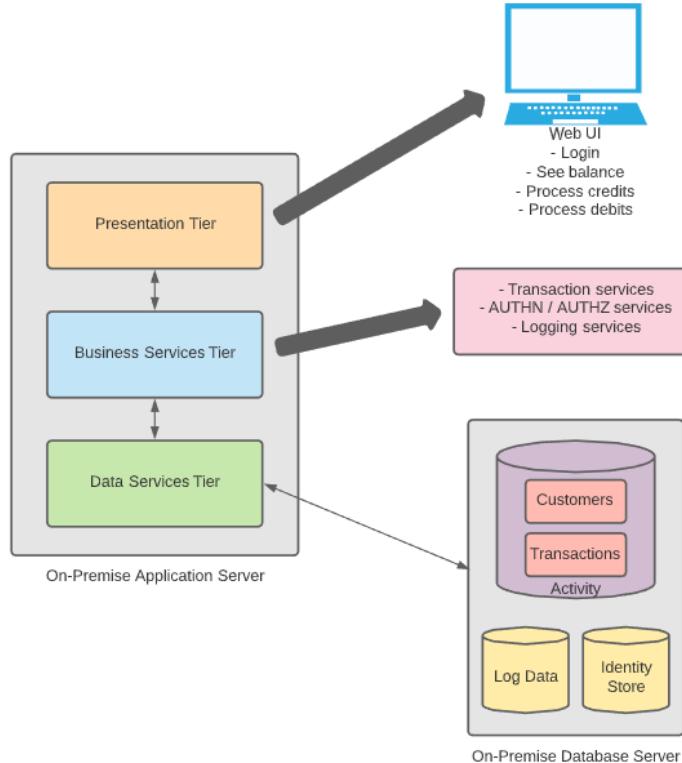
# The 6 R's – Application Migration Strategies

When moving from on-premise to the Cloud (or hybrid), the architect decides whether to migrate, modernize, or do some combination of the two



Source: <https://aws.amazon.com/blogs/enterprise-strategy/6-strategies-for-migrating-applications-to-the-cloud/>

# Current State Architecture

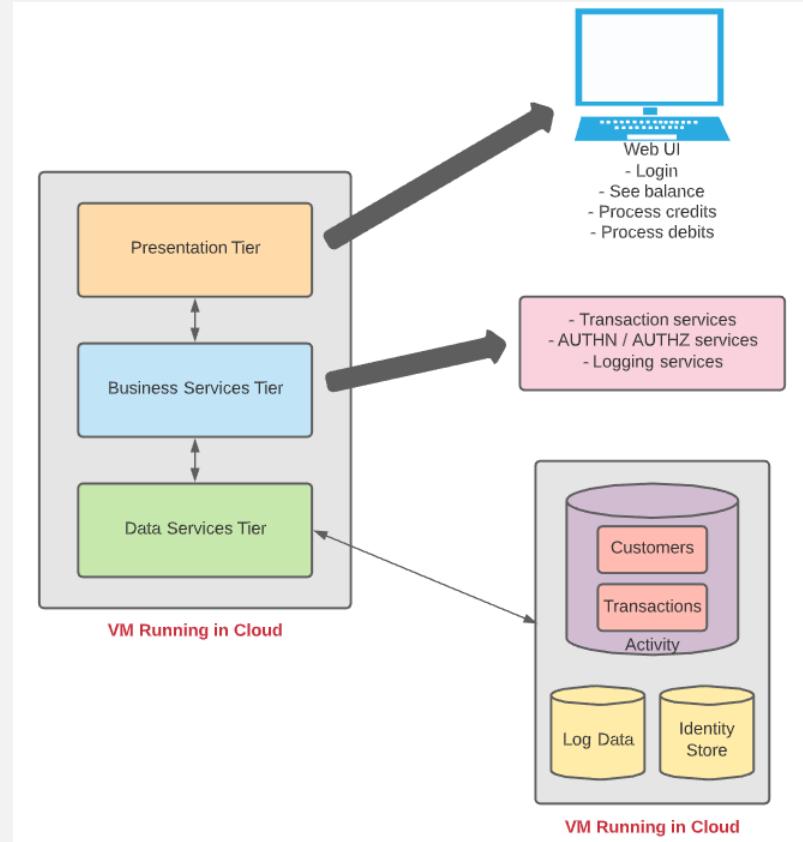




- AKA “lift & shift”
- For all intents and purposes, involves recreating the on-premise infrastructure in the Cloud
- Sometimes used to expedite retirement of a data center



- Can be a mechanism to quickly migrate workloads and see immediate cost savings, even without Cloud optimizations
- There are third-party tools available to help automate the migration
- Once the application is in the Cloud, it can be easier to apply Cloud optimizations vs. trying to migrate and optimize at the same time

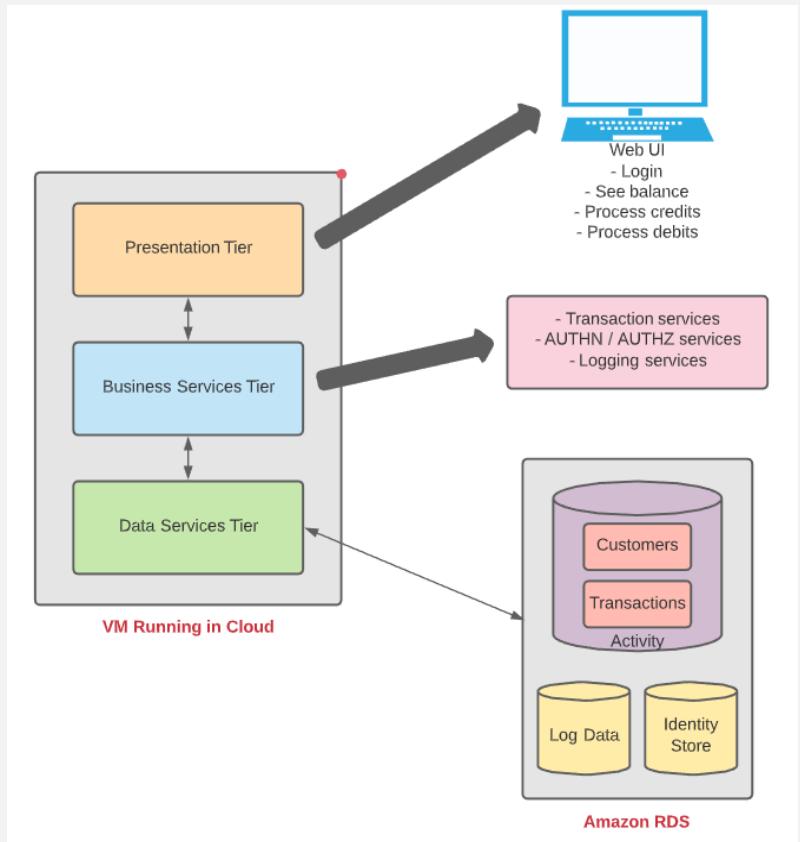




- AKA “lift, tinker & shift” or “lift & fiddle”
- Core architecture of the application will not change
- Involves recreating most of the on-premise infrastructure in the Cloud with a few Cloud optimizations



- Those optimizations will involve a move to one or more Cloud native services for a specific, tangible business benefit
- A common example is moving to a managed database (e.g., Relational Database Service, or RDS, in AWS)
- Enables migration speed while providing cost savings or benefit in a targeted portion of the application's architecture





- Involves rearchitecting the application to maximize utilization of Cloud native optimizations
- Likely the most expensive and most complex of the available options
- The application profile needs to fit, and business value must be identified commensurate with the cost required to execute



- Good option if the application can benefit from features, scalability, or performance offered by the Cloud
- Examples include rearchitecting a monolith to microservices running in the Cloud or moving an application to serverless technologies for scale (like the use case we've been discussing)



To Be Determined



- Good fit for applications that are candidates for moving from on-premise licensing (for installed products) to a SaaS model
- Often applications that have been installed to manage a specific type of business capability
- Examples can include Customer Relationship Management (CRM), Enterprise Resource Planning (ERP), HR or e-mail (among others)



- In some cases, an enterprise may have a “healthy” percentage of legacy applications that are no longer being used or maintained (especially for larger organizations)
- When completing the Application Portfolio Assessment and associated interviews with application owners, look for opportunities to recommend retire of the unused legacy apps
- This type of application can represent a “quick win”



- The associated savings can potentially be advantageously factored into the business case for the Cloud modernization effort
- Finally, retiring unused apps reduces attack surface area from a security perspective



- Applications that must remain in place but that cannot be migrated to the Cloud without major refactor
- This type of application profile may prevent the ability to completely move out of the data center (at least in the interim)
- From a cost perspective, it can be difficult for an enterprise to take on both the operating expense of Cloud while continuing to carry the capital expense of a data center and on-premise infrastructure
- The hybrid model can be a good solution to support where needed

# Reference Architectures

AWS Well-Architected Framework

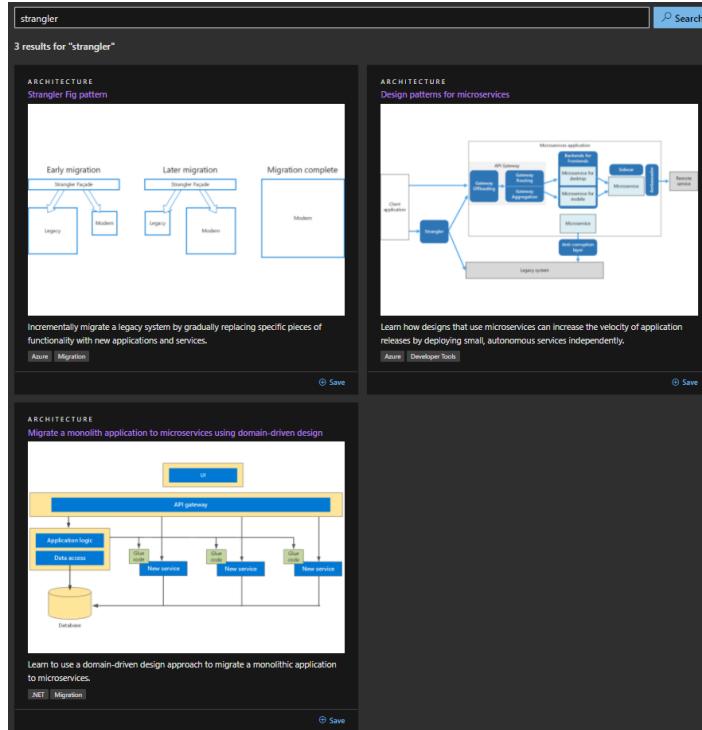
<https://aws.amazon.com/architecture/well-architected>

AWS Reference Architectures

<https://aws.amazon.com/architecture/>

# Azure Reference Architectures – Strangler Application

<https://docs.microsoft.com/en-us/azure/architecture/browse/?terms=strangler>



# Architecting for the Cloud



- The architect must choose between the various types of services provided by the CSP and design optimal integration between those services and components
- It may include a requirement to architect a solution that spans multiple providers or that links on-premise infrastructure & services with cloud-based resources
- Because we lack direct control over Cloud infrastructure and integrations are executed over the Internet, special care needs to be given to architectural concerns like security, resiliency, BC/DR (business continuity/data recovery), and performance
- Additionally, data classification and data sensitivity must be considered and accounted for – what infrastructure (in what areas of the world) is our data flowing through?
- What levels of dedicated connectivity and secrets management are required to ensure that we meet the requirements of our clients and any relevant regulators?

# Architecting for the Cloud



Hosting approach  
& cloud services  
to be utilized



Security  
requirements &  
effective  
method(s) to  
implement



Observability  
requirements &  
operational  
management  
strategy



Data  
management,  
classification, &  
protection



Quality of Service  
(QoS)  
requirements &  
the “-ilities”

## Sample Logical / Physical Architecture Diagrams

# Use Case – Architecting for the Cloud

As a squad, you would like to “sell” to the business a modernization effort to migrate a **monolithic, mainframe-based application** used to manage a **critical sales & distribution process** within the company to a **microservices** architecture. Included in the target goals would be the transition of the workflow from the data center to **Cloud and Cloud native technologies** (where possible).

Currently, the application leverages a mainframe-based relational data store to store and serve all data within the large and tightly-coupled workflow. There are several legacy technologies in play and the platform upon which the application is deployed is **several versions behind latest**. Key reports that the business needs for critical decisioning exist but **cannot be delivered in real-time**. They are requested in batch and sent as a follow-up, sometimes 1 or 2 days later.

Also, the business would like to be able to leverage the different dimensions of data housed within this system for **Machine Learning** purposes but, at present, the data is difficult to extract and utilize. The current mainframe system supports **external customer access through a Web portal**, **internal access for account & order management (including remote sales associates utilizing a mobile device)**, and a handful of API endpoints that are used by **external partners** to enable **authorized third-party sales**.

In this scenario, assume the intent is to migrate the application out of the data center and to the Cloud:

- What types of Cloud service (IaaS, PaaS, FaaS, SaaS) would you recommend leveraging and for which components?

# Architecture Governance



# The Case for Architecture Governance

- As previously mentioned, architecting modern systems requires us to balance competing concerns
- Every approach selected likely corresponds to a set of other options that were dismissed
- Understanding the “why” for a particular decision and vetting that “why” with other peers can help us move forward with confidence
- Additionally, there is value in formalizing a verified solution to a technical problem so that we can reuse that solution when a similar problem arises in the future
- We want to avoid having an enterprise of unicorns



# What Is It?

- A formalized process
- Should involve representatives from across the organization (ensure adequate understanding of all perspectives)
- Involves identifying, documenting, and cataloguing architectural standards and best practices
- When a team solves a complex architectural problem, we should codify so others can benefit



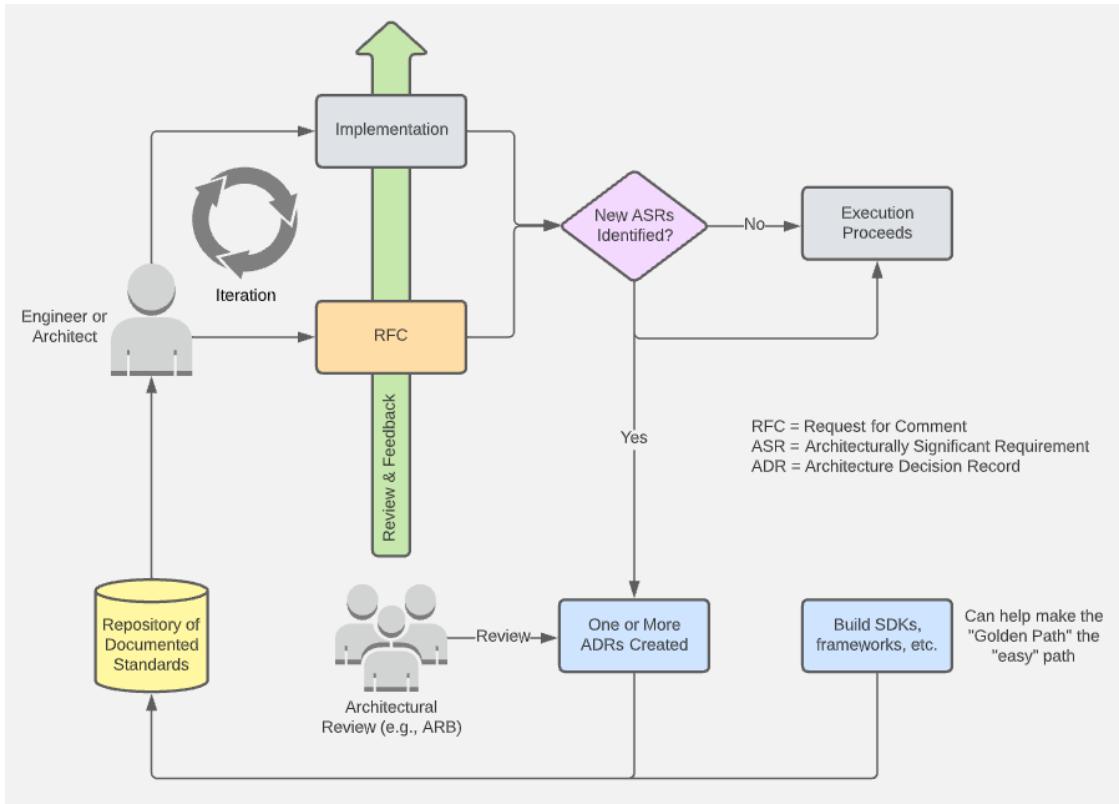
# What Is It?

- Engineers and architects looking to design software that meets a new business requirement can use this “repository” of standards & best practices for guidance
- Also, formalized reviews of an approach to solving an architecturally significant concern can help to “crowdsource” architectural quality

Think back to the discussion on “fitness functions” – architectural governance can be one of those “fitness functions” that helps guide our change



# What Does (Might) it Look Like?



# Keys to Remember

- Architecture Governance is not an “ivory tower” from which only select members of the organization decree “from on high”
- Requires a partnership between architecture, engineering, and business – architecture governance is about running architecture “as a business”
- Cross-organizational representation and input helps to ensure less gaps in understanding and a sense of “shared ownership” that can facilitate adherence



# Design Patterns

# Design Patterns



- As previously discussed, design patterns are proven solutions to a specific technical problem
- Focused primarily on the level of the source code
- Different classes of problem/solution that can be used (and reused) as building blocks

# Factory Pattern



- Creational pattern used to abstract creation logic from client
- Rather than create directly, code uses the factory to generate new instances
- Provides way to vary what gets created (and how) based on business logic

## LAB 09:

### Factory Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab09>

# Abstract Factory Pattern



- Creational pattern used to abstract creation logic from client
- Sometimes called factory of factories

## **LAB 10:**

Abstract Factory Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab10>

# Singleton Pattern



- Creational pattern used to manage class instance
- Seeks to create a single instance and provide a point of global access to that single instance
- Supports early and lazy instantiation patterns
- Likely for a specific type of use case – e.g., configuration detail

## **LAB 11:**

### Singleton Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab11>

# Prototype Pattern



- Creational pattern used to create instances by cloning rather than creating new
- Can be useful when intent is to minimize the number of instances or when cost of creating a new instance is high/complex

## **LAB 12:**

### Prototype Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab12>

# Builder Pattern



- Creational pattern used to build out an object over the course of multiple steps
- Each of the steps can be represented by a different object (if required)
- Enables clear codification of the construction approach

## **LAB 13:**

### Builder Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab13>

# Adapter Pattern



- Structural pattern used to wrap an interface
- Enables internal structure of object to remain as-is with layering in of a new class that can adapt specific client needs
- Allows reuse of existing functionality (even if not directly compatible)

## LAB 14:

### Adapter Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab14>

# Bridge Pattern



- Structural pattern used to decouple abstraction from implementation
- Supports extensibility and hides implementation details from code that references
- Allows abstraction and implementation to vary independently

## LAB 15:

### Bridge Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab15>

# Composite Pattern



- Structural pattern used to enable reference and access to a group of objects using a top-level object reference
- Can be used to represent a tree or object hierarchy
- Useful for cases like org charts or graphs of related objects

## **LAB 16:**

### Composite Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab16>

# Flyweight Pattern



- Structural pattern used to reduce the number of objects created, thereby reducing memory footprint and helping to improve performance
- Seeks to reuse objects where possible, only creating new if existing not found
- Supports shareable state through exposure of same object (if available) to multiple clients

## **LAB 19:**

### Flyweight Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab19>

# Decorator Pattern



- Structural pattern used to add new functionality to an existing object without altering its underlying, internal structure
- Clear example of OCP (Open for Extension, Closed for Modification)
- Uses composition instead of inheritance

## **LAB 17:**

### Decorator Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab17>

# Facade Pattern



- Structural pattern used to present a simplified interface that can be leveraged to facilitate access to a more complex, underlying structure
- Hides complexities, making downstream component/structure easier to use
- Promotes loose coupling through simplified abstraction

## **LAB 18:**

### Facade Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab18>

# Chain of Responsibility Pattern



- Behavioral pattern used to process a given request through a chain of objects, any of which can handle the request
- Reduces coupling by giving multiple classes the opportunity to service the request
- Enables a sequence of options, processing through the sequence until an appropriate handler is found

## **LAB 20:**

### Chain of Responsibility Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab20>

# Command Pattern



- Behavioral pattern used to encapsulate a handler for a specific type of command in a separate object
- Provides separation between object that issues the command and the agent fulfilling it
- Allows representation of a set of logic as a “thing” that can be passed around for encapsulated execution when needed

## **LAB 21:**

### Command Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab21>

# Observer Pattern



- Behavioral pattern used to manage a one-to-many relationship where multiple objects are watching and listening for activity in another
- Supports publish-subscribe approach – one publisher but can have multiple subscribers
- Event handling in programming languages is an example of the observer pattern

## **LAB 22:**

### Observer Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab22>

# State Pattern



- Behavioral pattern used to enable changes to a class's behavior based on its state and state transitions
- Manages state-specific behavior and keeps transitions explicit
- Uses a controller or context to drive behavior based on updates to state

## **LAB 23:**

### State Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab23>

# Strategy Pattern



- Behavioral pattern used to handle changes to class behavior or algorithm at run time in response to executed logic
- Uses a controller or context to select algorithm to be executed based on assigned strategy

## **LAB 24:**

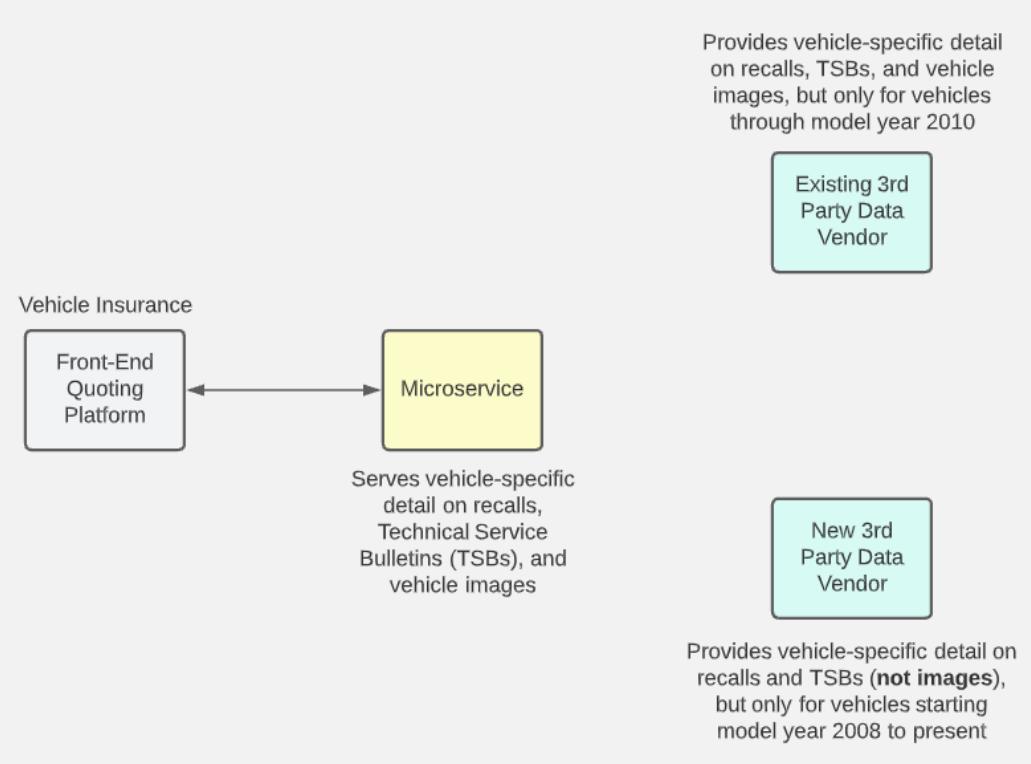
### Strategy Pattern

<https://github.com/KernelGamut32/advanced-software-architecture-public/tree/main/labs/lab24>

# Use Cases



# An Exercise in Flexibility



How would you architect a flexible solution (i.e., maintainable, loosely-coupled, etc.) to this business problem?

# SOLID in Practice

[https://github.com/KernelGamut32/advanced-software-architecture-public/blob/main/use-cases/Test%20Driven%20Development%20\(TDD\)%20-%20Class%20Project.pdf](https://github.com/KernelGamut32/advanced-software-architecture-public/blob/main/use-cases/Test%20Driven%20Development%20(TDD)%20-%20Class%20Project.pdf)

# Long Live the Mainframe

<https://github.com/KernelGamut32/advanced-software-architecture-public/blob/main/use-cases/Mainframe%20Arch.pdf>

# It's the Most Wonderful Time of the Year

<https://github.com/KernelGamut32/advanced-software-architecture-public/blob/main/hackathon/Toyshop.pdf>

## The Battle of Containers vs. Serverless

You have been tasked with building a system of microservices (APIs) to support a complex, end-to-end business workflow that involves several business domains. As a team, you are trying to decide between a couple of different implementation options – containers (Docker/Kubernetes) or serverless (AWS Lambda, Azure Functions).

As a group, discuss the merits of each and come prepared (upon return to the larger group) to talk about situations (and associated reasoning) where you would recommend one vs. the other. What are the pros & cons of each? Diagram a high-level reference architecture for the option (or options) you are recommending. Assume a target of ~15 microservices/APIs.

Don't forget to include consideration for important QoS (Quality of Service) requirements – scalability, reliability, performance, redundancy, maintainability, etc. Be prepared as a team to speak explicitly to those areas of concern.



# Thank you!



If you have additional questions,  
please reach out to me at:  
[asanders@gamuttechnologysvcs.com](mailto:asanders@gamuttechnologysvcs.com)